

443 Project

2022-11-06

Data Cleaning and Preprocessing

```
# read csv
df <- read.csv("Data_Group11.csv")
```

We first convert the data in to a more usable long format.

```
# convert month and day to MM-DD format
singledig <- which(df$X.1 < 10)
df$day <- as.character(df$X.1)
df$day[singledig] <- paste0("0", df$day[singledig])

df$month <- rep(c("01","02","03","04", "05","06","07","08","09","10","11","12"), times=c(31,29,31,30,31,30,31,31,30,31,30,31))

df$MMDD <- paste0(df$month, "-", df$day)

# drop unnecessary columns
to_drop <- c("X", "X.1", "X.2", "X1981.2010.mean", "X1981.2010.median", "month", "day")
df <- subset(df, select=!(names(df) %in% to_drop))

# melt years to be one column
library(tidyr)
df <- pivot_longer(data=df, cols=!MMDD, names_to="year", values_to="extent")

# format year and creating a column with YY-MM-DD format
library(stringr)
df$year = str_replace(df$year, "X", "")
df$YYMMDD <- paste0(df$year, "-", df$MMDD)

# tell R YYMMDD is a date
df$YYMMDD = as.Date(df$YYMMDD) # also conveniently rids of non leap year Feb 29's

# drop unnecessary columns
to_drop <- c("MMDD", "year")
df <- subset(df, select=!(names(df) %in% to_drop))

# order data by date
df <- df[order(df$YYMMDD),]

head(df)

## # A tibble: 6 x 2
##   extent YYMMDD
##   <dbl> <date>
## 1      NA 1978-01-01
```

```
## 2      NA 1978-01-02
## 3      NA 1978-01-03
## 4      NA 1978-01-04
## 5      NA 1978-01-05
## 6      NA 1978-01-06
```

We then deal with NA values.

```
# drop initial and ending NAs because we don't have data collected for these dates
library(zoo)
```

```
##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric

df <- na.trim(df)
```

There seems to be a change point at 1988-01-13 where a new method of measurement may of been put in to place. Before this date, there is a long stretch of NA values, and all measurements previously were recorded every second day.

```
# There is a period between 1987-12-03 and 1988-01-12 with a stretch of NAs. The following allows this ;
df_subset <- subset(df, YMMDD>as.Date("1987-12-01") & YMMDD<as.Date("1988-01-15"))
print(df_subset)
```

```
## # A tibble: 44 x 2
##   extent YMMDD
##   <dbl> <date>
## 1  12.6 1987-12-02
## 2    NA 1987-12-03
## 3    NA 1987-12-04
## 4    NA 1987-12-05
## 5    NA 1987-12-06
## 6    NA 1987-12-07
## 7    NA 1987-12-08
## 8    NA 1987-12-09
## 9    NA 1987-12-10
## 10   NA 1987-12-11
## # ... with 34 more rows
```

To deal with this long stretch of NAs and the NA values caused by measurement every second day, we propose two options. Either, we impute the missing data using cubic splines, or drop all observations before the possible changepoint.

```
# impute missing values using cubic splines as one option
df_imputed <- df
df_imputed$extent <- na.spline(df_imputed$extent)

# drop all data before 1988-01-13 as another option
df <- subset(df, YMMDD>=as.Date("1988-01-13"))
```

We are interested in overall trend, not day to day fluctuations, so we consider aggregating values by month.

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(lubridate)
```

```
##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union
```

```
df_aggregated <- df %>%
  group_by(year=year(YMMMDD), month=month(YMMMDD)) %>%
  mutate(avg_extent = mean(extent)) %>%
  distinct(year, month, .keep_all=TRUE) %>%
  subset(select=c(year, month, avg_extent))
```

Now that our data is cleaned and processed, we may proceed with analysis.

```
library(huxtable)

##
## Attaching package: 'huxtable'

## The following object is masked from 'package:dplyr':
##
##   add_rownames

summary_Avg_Extent <- summary(df_aggregated$avg_extent)
summary_Extent <- summary(df$extent)
summary <- as.data.frame(rbind(matrix(summary_Avg_Extent,nrow=1), matrix(summary_Extent,nrow=1)),
  row.names=c("Aggregated", "Unaggregated"))
colnames(summary) <- c("Min", "First Quartile", "Median", "Mean", "Third Quartile", "Max")

summary_table <- hux(summary, add_rownames = "")


summary_table %>%
  set_number_format(3) %>%
  set_align(everywhere, c(2,3,4,5,6,7), "center") %>%
  set_bottom_border(1, everywhere)
```

	Min	First Quartile	Median	Mean	Third Quartile	Max
Aggregated	3.566	8.276	11.914	11.150	14.166	15.957
Unaggregated	3.340	8.303	11.882	11.127	14.109	16.309

Aggregated Data


We analyze the monthly aggregated data.

```
# make a ts object
Avg_ExtentTS <- ts(df_aggregated$avg_extent, frequency=12, start=year(df$YYMMDD[1]))
plot(Avg_ExtentTS, ylab="Average Ice Extent", main="Average Ice Extent VS Time")
```




```
# Plotting Classical Decomposition
plot(decompose(Avg_ExtentTS)$trend, ylab="Average Ice Extent", main="Trend Component")
```

Trend Component




```
plot(decompose(Avg_ExtentTS)$season, ylab="Average Ice Extent", main="Seasonal Component")
```

Seasonal Component



```
plot(decompose(Avg_ExtentTS)$random, ylab="Average Ice Extent", main="Random Component")
```

Random Component



From the decomposition, we see there is a significant seasonal pattern, and likely significant trend.

Variance

From the plot of the data, we see a clear seasonal pattern, and perhaps a decreasing linear trend.

It is unclear whether variance is constant. We test this using the Fligner-Killeen test.

```
# do Fligner test for constant variance.
```

```
segments = factor(c(rep(1:4, each=84), rep(5, times=82)))
```

```
fligner.test(Avg_ExtentTS, segments)
```

```
##
```

```
## Fligner-Killeen test of homogeneity of variances
```

```
##
```

```
## data: Avg_ExtentTS and segments
```

```
## Fligner-Killeen:med chi-squared = 5.8565, df = 4, p-value = 0.2101
```

```
segments = factor(c(rep(1:9, each=42), rep(10, times=40)))
```

```
fligner.test(Avg_ExtentTS, segments)
```

```
##
```

```
## Fligner-Killeen test of homogeneity of variances
```

```
##
```

```
## data: Avg_ExtentTS and segments
```

```
## Fligner-Killeen:med chi-squared = 8.2771, df = 9, p-value = 0.5065
```

```
segments = factor(c(rep(1:19, each=21), rep(20, times=19)))
```

```
fligner.test(Avg_ExtentTS, segments)
```

```
##
```

```
## Fligner-Killeen test of homogeneity of variances
```

```
##
```

```
## data: Avg_ExtentTS and segments
## Fligner-Killeen:med chi-squared = 13.213, df = 19, p-value = 0.8275
segments = factor(c(rep(1:34, each=12), rep(35, times=10))) # corresponds to number of years of data
fligner.test(Avg_ExtentTS, segments)
```

```
##
## Fligner-Killeen test of homogeneity of variances
##
## data: Avg_ExtentTS and segments
## Fligner-Killeen:med chi-squared = 12.564, df = 34, p-value = 0.9997
```

All give high p-value so may conclude that variance is relatively constant. This is against expectation, but perhaps this is because the change in variance is not significant over such a small time frame.

```
# define mse function for future use
mse <- function(y, yhat) {
  return(mean((as.vector(y)-as.vector(yhat))^2))
}
```

First, split the data, in to train and test set.

```
Avg_ExtentTS_Train <- window(Avg_ExtentTS, 1988, 2020+11/12)
Avg_ExtentTS_Test <- window(Avg_ExtentTS, 2021, 2022+9/12)
```

Regression

Try to remove non-stationarity using Regression (Multiple Linear, Ridge, Lasso, Elastic Net).

Multiple Linear Regression

```
tim <- as.vector(time(Avg_ExtentTS_Train))
season <- factor(cycle(Avg_ExtentTS_Train))

# degree 1 polynomial of time
mlr_train <- lm(Avg_ExtentTS_Train~tim+season)

new <- data.frame(tim=as.vector(time(Avg_ExtentTS_Test)), season=factor(cycle(Avg_ExtentTS_Test)))
pmse_mlr <- mse(Avg_ExtentTS_Test, predict.lm(mlr_train, new))

# degree 2 polynomial of time
mlr_train_2 <- lm(Avg_ExtentTS_Train~poly(tim,2)+season)
pmse_mlr_2 <- mse(Avg_ExtentTS_Test, predict.lm(mlr_train_2, new))

# degree 3 polynomial of time
mlr_train_3 <- lm(Avg_ExtentTS_Train~poly(tim,3)+season)
pmse_mlr_3 <- mse(Avg_ExtentTS_Test, predict.lm(mlr_train_3, new))
```

The cubic model performs best on the hold out set.


```
# plot of cubic data
plot(Avg_ExtentTS)

# plot of fit
points(time(Avg_ExtentTS_Train), predict.lm(mlr_train_3), type='l', col='red')


# plot of test set prediction
```

```
points(time(Avg_ExtentTS_Test), predict.lm(mlr_train_3, new), type='l', col='blue')
```

```
# plot line at test set cutoff  
abline(v=2021, lty="dashed")
```




```
# plot of model residuals  
plot(mlr_train_3$residuals, type="l")
```



```
# plot of acf of residuals  
acf(mlr_train_3$residuals)
```


Series mlr_train_3\$residuals



Checking normality of this model:


```
# training the cubic model on the entire data
tim <- as.vector(time(Avg_ExtentTS))
season <- factor(cycle(Avg_ExtentTS))
mlr_3 <- lm(Avg_ExtentTS~poly(tim,3)+season)

# model diagnostics
par(mfrow=c(1,3)) # Dividing the plotting page into 4 panels
plot(mlr_3$fitted, mlr_3$residuals , pch=16 , col=adjustcolor("black" , 0.7), xlab="Fitted Values", ylab="Residuals")
title(main = "MLR With Polynomial Degree p=3")
abline(h=0,lty=2 , lwd=2 , col="red") # plotting a horizontal line at 0
car::qqPlot(mlr_3$residuals , pch=16, xlab="Theoretical Quantiles", ylab="Sample Quantiles")

## [1] 297 394

title(main = "MLR With Polynomial Degree p=3")
plot(mlr_3$residuals, pch=16 , col=adjustcolor("black" , 0.7), xlab="Time Index", ylab="Residuals") # p
title(main = "MLR With Polynomial Degree p=3")
abline(h=0,lty=2 , lwd=2 , col="red") # plotting a horizontal line at 0
```

MLR With Polynomial Degree $p=$ MLR With Polynomial Degree $p=$ MLR With Polynomial Degree $p=$



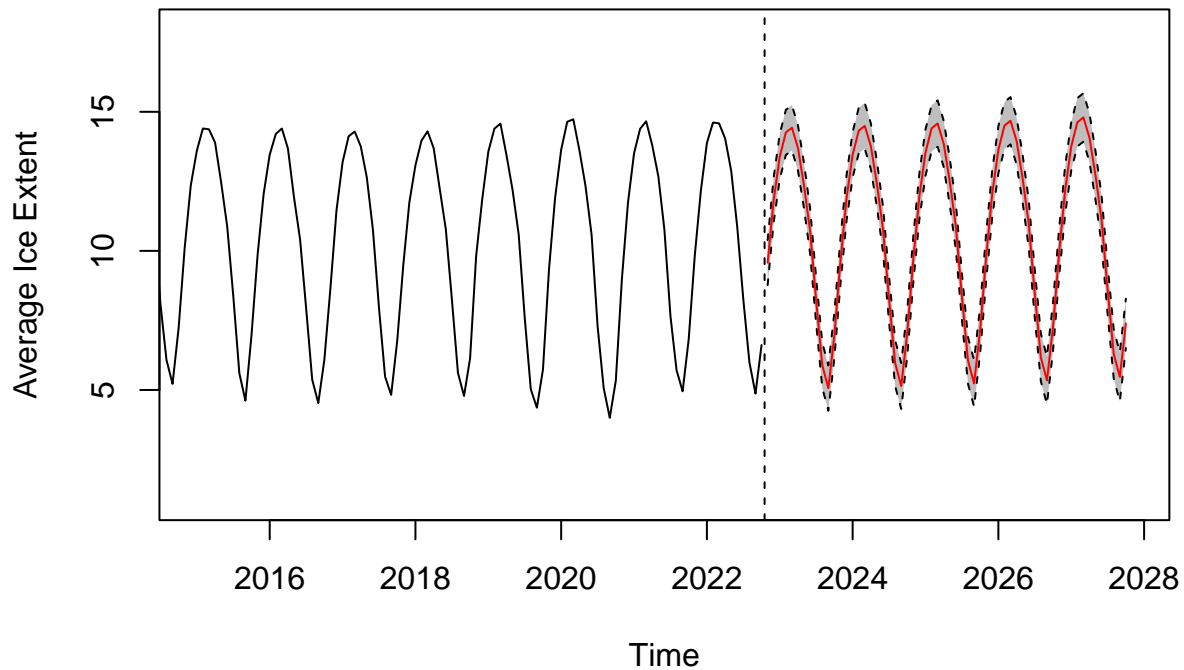
```
# Getting predictions for model and plotting
tim.new <- as.vector(seq(2022+10/12,2027+9/12,by=1/12))
season.new <- factor(c(11, 12, rep(1:12,4), 1:10))
new <- data.frame(tim=tim.new,season=season.new)
predict_mlr_3 <- predict.lm(mlr_3, new, interval='prediction')

par(mfrow=c(1,1))
plot(Avg_ExtentTS , xlim = c(2015 , 2027+10/12), ylim=c(1,18), ylab="Average Ice Extent", main="Predict.

#The three lines below plot the prediction interval in a grey scale
x = c(tim.new , rev(tim.new))
y = c(predict_mlr_3[, "upr"] , rev(predict_mlr_3[, "lwr"]))
polygon(x, y, col="grey", border=NA)

#The three line below add the predicted values and highlight the borders of the prediction interval
lines(x=tim.new, y=predict_mlr_3[, "upr"], col="black" , lty=2)
lines(x=tim.new, y=predict_mlr_3[, "lwr"], col="black", lty=2)
lines(x=tim.new, y=predict_mlr_3[, "fit"] , col="red")
abline(v=2022+9.5/12, lty="dashed")
```

Prediction from MLR Degree 3



```
# Shapiro-Wilk Test of normality test
shapiro.test(mlr_3$residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  mlr_3$residuals
## W = 0.98352, p-value = 0.0001077
```

Ridge

```
library(glmnet)
```

```
## Loading required package: Matrix
##
## Attaching package: 'Matrix'
## The following objects are masked from 'package:tidyr':
##
##   expand, pack, unpack
## Loaded glmnet 4.1-4
```

```
# fit to training data
tim <- as.vector(time(Avg_ExtentTS_Train))
season <- factor(cycle(Avg_ExtentTS_Train))
X <- model.matrix(as.vector(Avg_ExtentTS_Train)~tim+season)
X_2 <- model.matrix(as.vector(Avg_ExtentTS_Train)~poly(tim,2)+season)
X_3 <- model.matrix(as.vector(Avg_ExtentTS_Train)~poly(tim,3)+season)
ridge_train <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=0)
ridge_train_2 <- glmnet(X_2, as.vector(Avg_ExtentTS_Train), alpha=0)
```

```

ridge_train_3 <- glmnet(X_3, as.vector(Avg_ExtentTS_Train), alpha=0)

# compute mse on training data for each value of lambda
ridge_train_fitted <- predict(ridge_train, X)
ridge_train_fitted_2 <- predict(ridge_train_2, X_2)
ridge_train_fitted_3 <- predict(ridge_train_3, X_3)

mses <- c()
mses_2 <- c()
mses_3 <- c()
for(i in 1:100) {
  mses <- c(mses, mse(Avg_ExtentTS_Train, ridge_train_fitted[,i]))
  mses_2 <- c(mses_2, mse(Avg_ExtentTS_Train, ridge_train_fitted_2[,i]))
  mses_3 <- c(mses_3, mse(Avg_ExtentTS_Train, ridge_train_fitted_3[,i]))
}

min10_mses <- head(sort(mses), 10)
min10_mses_2 <- head(sort(mses_2), 10)
min10_mses_3 <- head(sort(mses_3), 10)

ridge_train_lambdas <- c()
ridge_train_lambdas_2 <- c()
ridge_train_lambdas_3 <- c()

for(m in min10_mses) {
  ridge_train_lambdas <- c(ridge_train_lambdas, ridge_train$lambda[which(mses==m)])
}
for(m in min10_mses_2) {
  ridge_train_lambdas_2 <- c(ridge_train_lambdas_2, ridge_train_2$lambda[which(mses_2==m)])
}
for(m in min10_mses_3) {
  ridge_train_lambdas_3 <- c(ridge_train_lambdas_3, ridge_train_3$lambda[which(mses_3==m)])
}

# retrain using lambdas that gave the 10 best fits
ridge_train <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=0, lambda=ridge_train_lambdas)
ridge_train_2 <- glmnet(X_2, as.vector(Avg_ExtentTS_Train), alpha=0, lambda=ridge_train_lambdas_2)
ridge_train_3 <- glmnet(X_3, as.vector(Avg_ExtentTS_Train), alpha=0, lambda=ridge_train_lambdas_3)

# predict the test set
tim <- as.vector(time(Avg_ExtentTS_Test))
season <- factor(cycle(Avg_ExtentTS_Test))
X <- model.matrix(as.vector(Avg_ExtentTS_Test)~tim+season)
X_2 <- model.matrix(as.vector(Avg_ExtentTS_Test)~poly(tim,2)+season)
X_3 <- model.matrix(as.vector(Avg_ExtentTS_Test)~poly(tim,3)+season)
ridge_predictions <- predict(ridge_train, X)
ridge_predictions_2 <- predict(ridge_train_2, X_2)
ridge_predictions_3 <- predict(ridge_train_3, X_3)

# compute pmse on test set
pmses <- c()
pmses_2 <- c()
pmses_3 <- c()
for(i in 1:10) {

```

```

pmses <- c(pmses, mse(Avg_ExtentTS_Test, ridge_predictions[,i]))
pmses_2 <- c(pmses_2, mse(Avg_ExtentTS_Test, ridge_predictions_2[,i]))
pmses_3 <- c(pmses_3, mse(Avg_ExtentTS_Test, ridge_predictions_3[,i]))
}

```

```

lambda_ridge <- ridge_train$lambda[which.min(pmses)]
pmse_ridge <- pmses[which.min(pmses)]
lambda_ridge_2 <- ridge_train_2$lambda[which.min(pmses_2)]
pmse_ridge_2 <- pmses_2[which.min(pmses_2)]
lambda_ridge_3 <- ridge_train_3$lambda[which.min(pmses_3)] #which.min
pmse_ridge_3 <- pmses_3[which.min(pmses_3)]

```

```
lambda_ridge
```

```
## [1] 0.1653056
```

```
pmse_ridge
```

```
## [1] 0.5345
```

```
lambda_ridge_2
```

```
## [1] 0.3818774
```

```
pmse_ridge_2
```

```
## [1] 5.203632
```

```
lambda_ridge_3
```

```
## [1] 0.3818774
```

```
pmse_ridge_3
```

```
## [1] 5.395475
```

We see that the degree 1 polynomial gives the lowest MSE.

```
#degree 1 polynomial
```

```
tim <- as.vector(time(Avg_ExtentTS_Train))
```

```
season <- factor(cycle(Avg_ExtentTS_Train))
```

```
X <- model.matrix(as.vector(Avg_ExtentTS_Train)~tim+season)
```

```
newX <- model.matrix(as.vector(Avg_ExtentTS_Test)~as.vector(time(Avg_ExtentTS_Test))+factor(cycle(Avg_E
```

```
ridge <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=0, lambda=0.1653056)
```


```
ridge_fitted <- predict(ridge, X)
```

```
plot(Avg_ExtentTS)
```


```
points(time(Avg_ExtentTS_Train), ridge_fitted, type='l', col='red')
```

```
points(time(Avg_ExtentTS_Test), predict(ridge, newX), type='l', col='blue')
```

```
abline(v=2021, lty="dashed")
```




```
ridge_residuals <- Avg_ExtentTS_Train - ridge_fitted  
plot(ridge_residuals, type="l")
```



```
acf(ridge_residuals)
```


s0




```
#degree 2 polynomial
tim <- as.vector(time(Avg_ExtentTS_Train))
season <- factor(cycle(Avg_ExtentTS_Train))
X <- model.matrix(as.vector(Avg_ExtentTS_Train)~poly(tim,2)+season)
newX <- model.matrix(as.vector(Avg_ExtentTS_Test)~poly(as.vector(time(Avg_ExtentTS_Test)),2)+factor(cycle(Avg_ExtentTS_Test))),2)+factor(cycle(Avg_ExtentTS_Test)))

ridge <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=0, lambda=0.3818774)
ridge_fitted <- predict(ridge, X)


plot(Avg_ExtentTS)
points(time(Avg_ExtentTS_Train), ridge_fitted, type='l', col='red')
points(time(Avg_ExtentTS_Test), predict(ridge, newX), type='l', col='blue')
abline(v=2021, lty="dashed")
```



```
ridge_residuals <- Avg_ExtentTS_Train - ridge_fitted
plot(ridge_residuals, type="l")
```



```
acf(ridge_residuals)
```

Lasso

```
library(glmnet)

# fit to training data
tim <- as.vector(time(Avg_ExtentTS_Train))
season <- factor(cycle(Avg_ExtentTS_Train))
X <- model.matrix(as.vector(Avg_ExtentTS_Train)~tim+season)
X_2 <- model.matrix(as.vector(Avg_ExtentTS_Train)~poly(tim,2)+season)
X_3 <- model.matrix(as.vector(Avg_ExtentTS_Train)~poly(tim,3)+season)
lasso_train <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=1)
lasso_train_2 <- glmnet(X_2, as.vector(Avg_ExtentTS_Train), alpha=1)
lasso_train_3 <- glmnet(X_3, as.vector(Avg_ExtentTS_Train), alpha=1)

# compute mse on training data for each value of lambda
lasso_train_fitted <- predict(lasso_train, X)
lasso_train_fitted_2 <- predict(lasso_train_2, X_2)
lasso_train_fitted_3 <- predict(lasso_train_3, X_3)

mses <- c()
mses_2 <- c()
mses_3 <- c()
for(i in 1:67) {
  mses <- c(mses, mse(Avg_ExtentTS_Train, lasso_train_fitted[,i]))
}
for(i in 1:68) {
  mses_2 <- c(mses_2, mse(Avg_ExtentTS_Train, lasso_train_fitted_2[,i]))
  mses_3 <- c(mses_3, mse(Avg_ExtentTS_Train, lasso_train_fitted_3[,i]))
}
```

```

}

min10_mses <- head(sort(mses), 10)
min10_mses_2 <- head(sort(mses_2), 10)
min10_mses_3 <- head(sort(mses_3), 10)

lasso_train_lambdas <- c()
lasso_train_lambdas_2 <- c()
lasso_train_lambdas_3 <- c()

for(m in min10_mses) {
  lasso_train_lambdas <- c(lasso_train_lambdas, lasso_train$lambda[which(mses==m)])
}
for(m in min10_mses_2) {
  lasso_train_lambdas_2 <- c(lasso_train_lambdas_2, lasso_train_2$lambda[which(mses_2==m)])
}
for(m in min10_mses_3) {
  lasso_train_lambdas_3 <- c(lasso_train_lambdas_3, lasso_train_3$lambda[which(mses_3==m)])
}

# retrain using lambdas that gave the 10 best fits
lasso_train <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=0, lambda=lasso_train_lambdas)
lasso_train_2 <- glmnet(X_2, as.vector(Avg_ExtentTS_Train), alpha=0, lambda=lasso_train_lambdas_2)
lasso_train_3 <- glmnet(X_3, as.vector(Avg_ExtentTS_Train), alpha=0, lambda=lasso_train_lambdas_3)

# predict the test set
tim <- as.vector(time(Avg_ExtentTS_Test))
season <- factor(cycle(Avg_ExtentTS_Test))
X <- model.matrix(as.vector(Avg_ExtentTS_Test)~tim+season)
X_2 <- model.matrix(as.vector(Avg_ExtentTS_Test)~poly(tim,2)+season)
X_3 <- model.matrix(as.vector(Avg_ExtentTS_Test)~poly(tim,3)+season)
lasso_predictions <- predict(lasso_train, X)
lasso_predictions_2 <- predict(lasso_train_2, X_2)
lasso_predictions_3 <- predict(lasso_train_3, X_3)

# compute pmse on test set
pmse <- c()
pmse_2 <- c()
pmse_3 <- c()
for(i in 1:10) {
  pmse <- c(pmse, mse(Avg_ExtentTS_Test, lasso_predictions[,i]))
  pmse_2 <- c(pmse_2, mse(Avg_ExtentTS_Test, lasso_predictions_2[,i]))
  pmse_3 <- c(pmse_3, mse(Avg_ExtentTS_Test, lasso_predictions_3[,i]))
}

lambda_lasso <- lasso_train$lambda[which.min(pmse)]
pmse_lasso <- pmse[which.min(pmse)]
lambda_lasso_2 <- lasso_train_2$lambda[which.min(pmse_2)]
pmse_lasso_2 <- pmse_2[which.min(pmse_2)]
lambda_lasso_3 <- lasso_train_3$lambda[which.min(pmse_3)]
pmse_lasso_3 <- pmse_3[which.min(pmse_3)]

lambda_lasso

```

```
## [1] 0.003561401
```

```
pmse_lasso
```

```
## [1] 0.368189
```

```
lambda_lasso_2
```

```
## [1] 0.007496408
```

```
pmse_lasso_2
```

```
## [1] 6.69867
```

```
lambda_lasso_3
```

```
## [1] 0.007496408
```

```
pmse_lasso_3
```

```
## [1] 6.85394
```

We see that the degree 1 polynomial gives the lowest MSE.

```
#degree 1 polynomial
```

```
tim <- as.vector(time(Avg_ExtentTS_Train))
```

```
season <- factor(cycle(Avg_ExtentTS_Train))
```

```
X <- model.matrix(as.vector(Avg_ExtentTS_Train)~tim+season)
```

```
newX <- model.matrix(as.vector(Avg_ExtentTS_Test)~as.vector(time(Avg_ExtentTS_Test))+factor(cycle(Avg_E
```

```
lasso <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=1, lambda=0.003561401)
```


```
lasso_fitted <- predict(lasso, X)
```

```
plot(Avg_ExtentTS)
```


```
points(time(Avg_ExtentTS_Train), lasso_fitted, type='l', col='red')
```

```
points(time(Avg_ExtentTS_Test), predict(lasso, newX), type='l', col='blue')
```


```
abline(v=2021, lty="dashed")
```



```
lasso_residuals <- Avg_ExtentTS_Train - lasso_fitted  
plot(lasso_residuals, type="l")
```



```
acf(lasso_residuals)
```



Elastic Net

```
library(glmnet)
alpha_seq <- seq(0.1, 0.9, by=0.1)

en_train_min_lamdass <- c()
en_train_min_lamdass_2 <- c()
en_train_min_lamdass_3 <- c()
min_pmses <- c()
min_pmses_2 <- c()
min_pmses_3 <- c()

for(a in alpha_seq){
  # fit to training data
  tim <- as.vector(time(Avg_ExtentTS_Train))
  season <- factor(cycle(Avg_ExtentTS_Train))
  X <- model.matrix(as.vector(Avg_ExtentTS_Train)~tim+season)
  X_2 <- model.matrix(as.vector(Avg_ExtentTS_Train)~poly(tim,2)+season)
  X_3 <- model.matrix(as.vector(Avg_ExtentTS_Train)~poly(tim,3)+season)
  en_train <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=a)
  en_train_2 <- glmnet(X_2, as.vector(Avg_ExtentTS_Train), alpha=a)
  en_train_3 <- glmnet(X_3, as.vector(Avg_ExtentTS_Train), alpha=a)

  # compute mse on training data for each value of lambda
  en_train_fitted <- predict(en_train, X)
  en_train_fitted_2 <- predict(en_train_2, X_2)
  en_train_fitted_3 <- predict(en_train_3, X_3)

  mses <- c()
  mses_2 <- c()
  mses_3 <- c()
  for(i in 1:length(en_train$lambda)) {
    mses <- c(mses, mse(Avg_ExtentTS_Train, en_train_fitted[,i]))
  }
  for(i in 1:length(en_train_2$lambda)) {
    mses_2 <- c(mses_2, mse(Avg_ExtentTS_Train, en_train_fitted_2[,i]))
  }
  for(i in 1:length(en_train_3$lambda)) {
    mses_3 <- c(mses_3, mse(Avg_ExtentTS_Train, en_train_fitted_3[,i]))
  }

  min10_mses <- head(sort(mses), 10)
  min10_mses_2 <- head(sort(mses_2), 10)
  min10_mses_3 <- head(sort(mses_3), 10)

  en_train_lamdass <- c()
  en_train_lamdass_2 <- c()
  en_train_lamdass_3 <- c()

  for(m in min10_mses) {
    en_train_lamdass <- c(en_train_lamdass, en_train$lambda[which(mses==m)])
  }
  for(m in min10_mses_2) {
    en_train_lamdass_2 <- c(en_train_lamdass_2, en_train_2$lambda[which(mses_2==m)])
  }
}
```

```

}
for(m in min10_mses_3) {
  en_train_lambdas_3 <- c(en_train_lambdas_3, en_train_3$lambda[which(mses_3==m)])
}

# retrain using lambdas that gave the 10 best fits
en_train <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=a, lambda=en_train_lambdas)
en_train_2 <- glmnet(X_2, as.vector(Avg_ExtentTS_Train), alpha=a, lambda=en_train_lambdas_2)
en_train_3 <- glmnet(X_3, as.vector(Avg_ExtentTS_Train), alpha=a, lambda=en_train_lambdas_3)

# predict the test set
tim <- as.vector(time(Avg_ExtentTS_Test))
season <- factor(cycle(Avg_ExtentTS_Test))
X <- model.matrix(as.vector(Avg_ExtentTS_Test)~tim+season)
X_2 <- model.matrix(as.vector(Avg_ExtentTS_Test)~poly(tim,2)+season)
X_3 <- model.matrix(as.vector(Avg_ExtentTS_Test)~poly(tim,3)+season)
en_predictions <- predict(en_train, X)
en_predictions_2 <- predict(en_train_2, X_2)
en_predictions_3 <- predict(en_train_3, X_3)

# compute pmse on test set
pmses <- c()
pmses_2 <- c()
pmses_3 <- c()
for(i in 1:10) {
  pmses <- c(pmses, mse(Avg_ExtentTS_Test, en_predictions[,i]))
  pmses_2 <- c(pmses_2, mse(Avg_ExtentTS_Test, en_predictions_2[,i]))
  pmses_3 <- c(pmses_3, mse(Avg_ExtentTS_Test, en_predictions_3[,i]))
}

min_pmses <- c(min_pmses, pmses[which.min(pmses)])
min_pmses_2 <- c(min_pmses_2, pmses_2[which.min(pmses_2)])
min_pmses_3 <- c(min_pmses_3, pmses_3[which.min(pmses_3)])

en_train_min_lambdas <- c(en_train_min_lambdas, en_train$lambda[which.min(pmses)])
en_train_min_lambdas_2 <- c(en_train_min_lambdas_2, en_train_2$lambda[which.min(pmses_2)])
en_train_min_lambdas_3 <- c(en_train_min_lambdas_3, en_train_3$lambda[which.min(pmses_3)])
}

pmse_en <- min_pmses[which.min(min_pmses)]
alpha_en <- alpha_seq[which.min(min_pmses)]
lambda_en <- en_train_min_lambdas[which.min(min_pmses)]
pmse_en_2 <- min_pmses_2[which.min(min_pmses_2)]
lambda_en_2 <- en_train_min_lambdas_2[which.min(min_pmses_2)]
alpha_en_2 <- alpha_seq[which.min(min_pmses_2)]
pmse_en_3 <- min_pmses_3[which.min(min_pmses_3)]
lambda_en_3 <- en_train_min_lambdas_3[which.min(min_pmses_3)]
alpha_en_3 <- alpha_seq[which.min(min_pmses_3)]

pmse_en

```

```
## [1] 0.3659694
```

```
lambda_en
```

```
## [1] 0.004342926
```

```
alpha_en
```

```
## [1] 0.9
```

```
pmse_en_2
```

```
## [1] 6.537295
```

```
lambda_en_2
```

```
## [1] 0.008329342
```

```
alpha_en_2
```

```
## [1] 0.9
```

```
pmse_en_3
```

```
## [1] 6.668906
```

```
lambda_en_3
```

```
## [1] 0.008329342
```

```
alpha_en_3
```

```
## [1] 0.9
```

We see that the degree 1 polynomial gives the lowest MSE.

```
#degree 1 polynomial
```

```
tim <- as.vector(time(Avg_ExtentTS_Train))
```

```
season <- factor(cycle(Avg_ExtentTS_Train))
```

```
X <- model.matrix(as.vector(Avg_ExtentTS_Train)~tim+season)
```

```
newX <- model.matrix(as.vector(Avg_ExtentTS_Test)~as.vector(time(Avg_ExtentTS_Test))+factor(cycle(Avg_E
```

```
en <- glmnet(X, as.vector(Avg_ExtentTS_Train), alpha=0.9, lambda=0.004342926)
```


```
en_fitted <- predict(en, X)
```

```
plot(Avg_ExtentTS)
```


```
points(time(Avg_ExtentTS_Train), en_fitted, type='l', col='red')
```

```
points(time(Avg_ExtentTS_Test), predict(en, newX), type='l', col='blue')
```

```
abline(v=2021, lty="dashed")
```




```
en_residuals <- Avg_ExtentTS_Train - en_fitted  
plot(en_residuals, type="l")
```



```
acf(en_residuals)
```


s0



```
# creating a table to summarise regression results
pmse <- c(pmse_mlr, pmse_mlr_2, pmse_mlr_3, pmse_ridge, pmse_ridge_2, pmse_ridge_3, pmse_lasso, pmse_lasso_2, pmse_lasso_3)
time_degree <- rep(c(1,2,3), 4)
lambda <- c("", "", "", lambda_ridge, lambda_ridge_2, lambda_ridge_3, lambda_lasso, lambda_lasso_2, lambda_lasso_3)
alpha <- c("", "", "", "", "", "", "", "", "", alpha_en, alpha_en_2, alpha_en_3)
regression <- as.data.frame(matrix(c(pmse, time_degree, lambda, alpha), ncol=4),
                                row.names=c("", "MLR", "", "", "Ridge", "", "", "Lasso", "", "", "Elastic Net"))
colnames(regression) <- c("Prediction MSE", "Degree of Time Polynomial", "Lambda", "Alpha")

regression_table <- hux(regression, add_rownames = "")

regression_table %>%
  set_number_format(3) %>%
  set_align(everywhere, everywhere, "center") %>%
  set_bottom_border(c(1,4,7,10), everywhere)
```

Holt-Winters

Try to remove non-stationarity using exponential smoothing, double exponential smoothing, additive HW, and multiplicative HW.

Exponential Smoothing

```
es <- HoltWinters(Avg_ExtentTS_Train, gamma = FALSE, beta = FALSE)
predict_es = predict(es, n.ahead=22)
pmse_es <- mse(Avg_ExtentTS_Test, predict_es)
```

	Prediction MSE	Degree of Time Polynomial	Lambda	Alpha
X	0.366	1.000		
MLR	0.455	2.000		
X0.100	0.233	3.000		
X0.200	0.535	1.000	0.165	
Ridge	5.204	2.000	0.382	
X0.300	5.395	3.000	0.382	
X0.400	0.368	1.000	0.004	
Lasso	6.699	2.000	0.007	
X0.500	6.854	3.000	0.007	
X0.600	0.366	1.000	0.004	0.900
Elastic.Net	6.537	2.000	0.008	0.900
X0.700	6.669	3.000	0.008	0.900

Double Exponential Smoothing

```
des <- HoltWinters(Avg_ExtentTS_Train, gamma = FALSE)
predict_des = predict(des, n.ahead=22)
pmse_des <- mse(Avg_ExtentTS_Test, predict_des)
```

No Trend

```
no_trend <- HoltWinters(Avg_ExtentTS_Train, beta = FALSE)
predict_no_trend = predict(no_trend, n.ahead=22)
pmse_no_trend <- mse(Avg_ExtentTS_Test, predict_no_trend)
```

Additive Holt-Winters

```
additive <- HoltWinters(Avg_ExtentTS_Train, seasonal = "additive")
predict_additive = predict(additive, n.ahead=22)
pmse_additive <- mse(Avg_ExtentTS_Test, predict_additive)
```

Multiplicative Holt-Winters

```
multiplicative <- HoltWinters(Avg_ExtentTS_Train, seasonal = "multiplicative")
predict_multiplicative = predict(multiplicative, n.ahead=22)
pmse_multiplicative <- mse(Avg_ExtentTS_Test, predict_additive)
```

```
pmse <- c(pmse_es, pmse_des, pmse_no_trend, pmse_additive, pmse_multiplicative)
HW <- as.data.frame(matrix(pmse, nrow=1),
```

```
row.names=c("Prediction MSE"))
colnames(HW) <- c("\nExponential\nSmoothing", "Double\nExponential\nSmoothing", "HW\nWithout\nTrend", "
```


```
HW_table <- hux(HW, add_rownames = "")

HW_table %>%
  set_number_format(3) %>%
  set_align(everywhere, everywhere, "center") %>%
  set_bottom_border(1, everywhere)
```

	ExponentialSmoothing	DoubleExponentialSmoothing	HWWithoutTrend	AdditiveHW	MultiplicativeHW
Prediction MSE	13.561	14.468	0.459	0.778	0.778


Best of HW models seems to be model with no trend. We fit this model to the entire data.

```
residuals_HW <- as.vector(Avg_ExtentTS_Train[which(time(Avg_ExtentTS_Train)>=1989)]) - no_trend$fitted[
plot(residuals_HW, type="l")
```




```
acf(residuals_HW, lag.max=36)
```

Series residuals_HW




```
predict_no_trend <- predict(no_trend, n.ahead=22, prediction.interval = TRUE , level=0.95)  
plot(no_trend, predict_no_trend)
```

Holt-Winters filtering




```
acf(residuals_HW, lag.max=36)
```

Series residuals_HW



```
pacf(residuals_HW, lag.max=36)
```

Series residuals_HW



Note here that a more thorough analysis might have applied Box-Jenkins on the HW residuals, because a case could be made that they are stationary.

Prediction With best HW:


```
# Getting predictions for HW and plotting
HW <- HoltWinters(Avg_ExtentTS, beta = FALSE)
predict_HW <- predict(HW, n.ahead=60, prediction.interval = TRUE , level=0.95)

plot(Avg_ExtentTS , xlim = c(2015 , 2027+10/12), ylim=c(1,18), ylab="Average Ice Extent", main="Predict.

#The three lines below plot the prediction interval in a grey scale
x = c(time(predict_HW[, "upr"]) , rev(time(predict_HW[, "upr"])))
y = c(predict_HW[, "upr"] , rev(predict_HW[, "lwr"]))
polygon(x, y, col="grey", border=NA)

#The three line below add the predicted values and highlight the borders of the prediction interval
lines(predict_HW[, "upr"], col="black" , lty=2)
lines(predict_HW[, "lwr"], col="black", lty=2)
lines(predict_HW[, "fit"] , col="red")
abline(v=2022+9.5/12, lty="dashed")
```

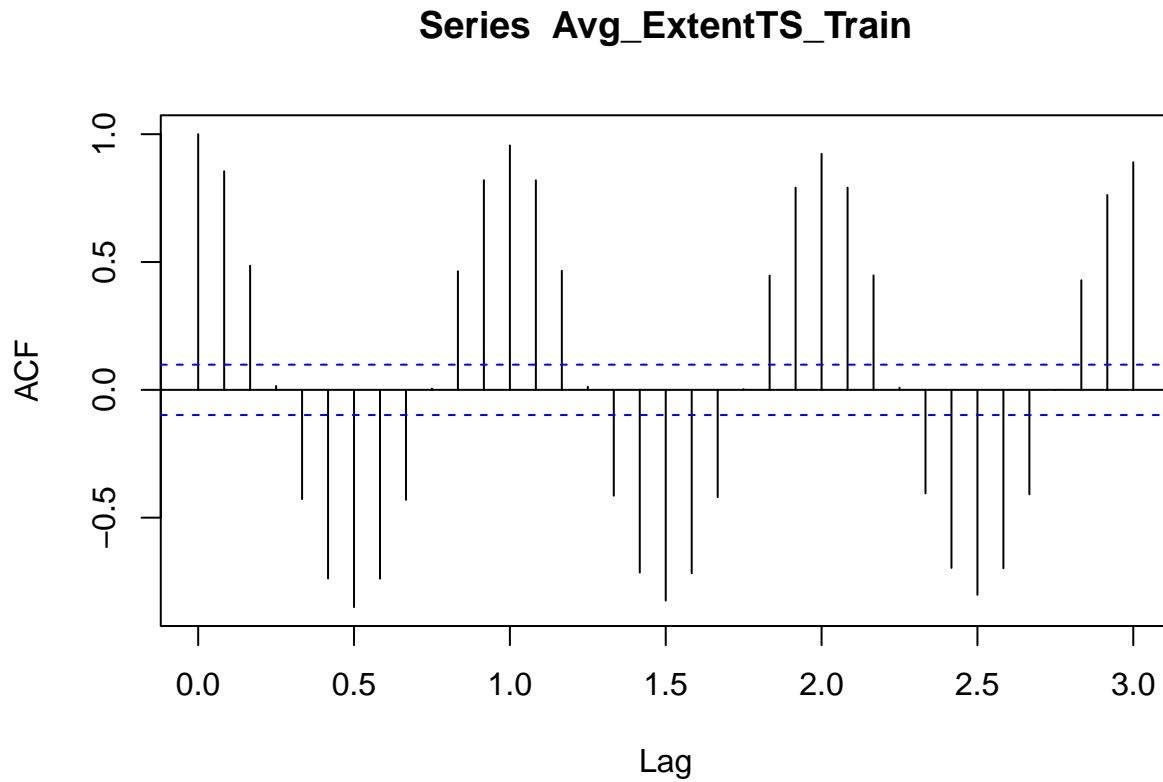
Prediction from Holt–Winters




Differencing on Train Data

Try differencing to remove non-stationarity.

```
acf(Avg_ExtentTS_Train, lag.max=36)
```




```
plot(Avg_ExtentTS_Train)
```




```
#differencing in lag of season  
diff12.Extent=diff(Avg_ExtentTS_Train, lag=12)  
acf(diff12.Extent, lag.max=50)
```

Series diff12.Extent




```
pacf(diff12.Extent, lag.max=50)
```

Series diff12.Extent




```
plot(diff12.Extent)
```





```
#regular differencing
diff.Extent=diff(Avg_ExtentTS_Train)
acf(diff.Extent, lag.max=36)
```

Series diff.Extent




```
plot(diff.Extent)
```




```
# seasonal+regular differencing
diff12.diff.Extent=diff(diff12.Extent)
acf(diff12.diff.Extent, lag.max=36)
```

Series diff12.diff.Extent




```
plot(diff12.diff.Extent)
```




It seems that regression performs poorly in terms of removing non-stationarity compared to HW and Differencing. HW and differencing seem to perform similarly. For simplicity, we proceed with differencing.
 #TODO reword this

```
acf(diff12.diff.Extent, lag.max=50)
```

Series diff12.diff.Extent



```
pacf(diff12.diff.Extent, lag.max=50)
```



We propose the following models:

Box-Jenkins on Seasonally differenced data

Seasonal differencing of the data seems to be enough to achieve stationarity, so we proceed with that. See appendix for analysis of seasonal + regular differencing


```
library(astsa)
```

```
#SARIMA(1,0,1)x(0,1,1)_12
```

```
model_21_train <- sarima(Avg_ExtentTS_Train, p=1, d=0, q=1, P=0, D=1, Q=1, S=12, details = TRUE)
```


```
## initial value -0.801086
## iter 2 value -1.133503
## iter 3 value -1.361046
## iter 4 value -1.364659
## iter 5 value -1.367122
## iter 6 value -1.367636
## iter 7 value -1.367710
## iter 8 value -1.367744
## iter 9 value -1.367751
## iter 10 value -1.367751
## iter 10 value -1.367751
## final value -1.367751
## converged
## initial value -1.365232
## iter 2 value -1.365266
## iter 3 value -1.365293
```

```
## iter 4 value -1.365301
## iter 5 value -1.365301
## iter 6 value -1.365301
## iter 6 value -1.365301
## iter 6 value -1.365301
## final value -1.365301
## converged
```



```
model_21_train_residuals = resid(model_21_train$fit)
hist(model_21_train_residuals)
```

Histogram of model_21_train_residuals



```
shapiro.test(model_21_train_residuals)
```


```
##  
##  Shapiro-Wilk normality test  
##  
## data:  model_21_train_residuals  
## W = 0.98802, p-value = 0.002432
```

```
#SARIMA(1,0,1)x(1,1,0)_12
```

```
model_22_train <- sarima(Avg_ExtentTS_Train, p=1, d=0, q=1, P=1, D=1, Q=0, S=12, details = TRUE)
```


```
## initial  value -0.794710  
## iter    2 value -1.125638  
## iter    3 value -1.297146  
## iter    4 value -1.306046  
## iter    5 value -1.306798  
## iter    6 value -1.306962  
## iter    7 value -1.307040  
## iter    8 value -1.307043  
## iter    9 value -1.307044  
## iter   10 value -1.307044  
## iter   11 value -1.307044  
## iter   12 value -1.307044  
## iter   12 value -1.307044  
## iter   12 value -1.307044  
## final    value -1.307044  
## converged  
## initial  value -1.307068  
## iter    2 value -1.307118  
## iter    3 value -1.307150
```

```
## iter 4 value -1.307158
## iter 5 value -1.307161
## iter 6 value -1.307161
## iter 7 value -1.307161
## iter 8 value -1.307161
## iter 8 value -1.307161
## iter 8 value -1.307161
## final value -1.307161
## converged
```



```
model_22_train_residuals = resid(model_22_train$fit)
hist(model_22_train_residuals)
```

Histogram of model_22_train_residuals



```
shapiro.test(model_22_train_residuals)
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_22_train_residuals  
## W = 0.98428, p-value = 0.0002677
```


```
#SARIMA(1,0,1)x(0,1,3)_12
```

```
model_23_train <- sarima(Avg_ExtentTS_Train, p=1, d=0, q=1, P=0, D=1, Q=3, S=12 , details = TRUE)
```

```
## initial  value -0.801086  
## iter    2 value -1.144412  
## iter    3 value -1.344981  
## iter    4 value -1.357046  
## iter    5 value -1.368996  
## iter    6 value -1.371115  
## iter    7 value -1.371624  
## iter    8 value -1.371650  
## iter    9 value -1.371662  
## iter   10 value -1.371662  
## iter   11 value -1.371662  
## iter   11 value -1.371662  
## iter   11 value -1.371662  
## final   value -1.371662  
## converged  
## initial  value -1.368723  
## iter    2 value -1.368752  
## iter    3 value -1.368772  
## iter    4 value -1.368806
```




```
## iter 5 value -1.368811
## iter 6 value -1.368813
## iter 7 value -1.368813
## iter 7 value -1.368813
## iter 7 value -1.368813
## final value -1.368813
## converged
```



```
model_23_train_residuals = resid(model_23_train$fit)
hist(model_23_train_residuals)
```

Histogram of model_23_train_residuals



```
shapiro.test(model_23_train_residuals)
```


```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_23_train_residuals  
## W = 0.98706, p-value = 0.001346
```

```
#SARIMA(1,0,1)x(0,1,4)_12
```

```
model_24_train <- sarima(Avg_ExtentTS_Train, p=1, d=0, q=1, P=0, D=1, Q=4, S=12, details = TRUE)
```


```
## initial value -0.801086  
## iter 2 value -1.140892  
## iter 3 value -1.368903  
## iter 4 value -1.371731  
## iter 5 value -1.373468  
## iter 6 value -1.375135  
## iter 7 value -1.375394  
## iter 8 value -1.375462  
## iter 9 value -1.375466  
## iter 10 value -1.375469  
## iter 11 value -1.375469  
## iter 11 value -1.375469  
## iter 11 value -1.375469  
## final value -1.375469  
## converged  
## initial value -1.372802  
## iter 2 value -1.372833  
## iter 3 value -1.372871  
## iter 4 value -1.372893
```

```
## iter 5 value -1.372905
## iter 6 value -1.372905
## iter 6 value -1.372905
## iter 6 value -1.372905
## final value -1.372905
## converged
```



```
model_24_train_residuals = resid(model_24_train$fit)
hist(model_24_train_residuals)
```

Histogram of model_24_train_residuals



```
shapiro.test(model_24_train_residuals)
```


```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_24_train_residuals  
## W = 0.98849, p-value = 0.003246
```

```
#SARIMA(1,0,1)x(1,1,1)_12
```

```
model_25_train <- sarima(Avg_ExtentTS_Train, p=1, d=0, q=1, P=1, D=1, Q=1, S=12, details = TRUE)
```


```
## initial value -0.794710  
## iter 2 value -1.152892  
## iter 3 value -1.325426  
## iter 4 value -1.345857  
## iter 5 value -1.352890  
## iter 6 value -1.359234  
## iter 7 value -1.360827  
## iter 8 value -1.361190  
## iter 9 value -1.361198  
## iter 10 value -1.361203  
## iter 11 value -1.361205  
## iter 12 value -1.361205  
## iter 12 value -1.361205  
## final value -1.361205  
## converged  
## initial value -1.364327  
## iter 2 value -1.364958  
## iter 3 value -1.365484  
## iter 4 value -1.365574
```

```
## iter 5 value -1.365589
## iter 6 value -1.365590
## iter 7 value -1.365590
## iter 8 value -1.365590
## iter 8 value -1.365590
## iter 8 value -1.365590
## final value -1.365590
## converged
```



```
model_25_train_residuals = resid(model_25_train$fit)
hist(model_25_train_residuals)
```

Histogram of model_25_train_residuals



```
shapiro.test(model_25_train_residuals)
```


```
##  
##  Shapiro-Wilk normality test  
##  
## data:  model_25_train_residuals  
## W = 0.98841, p-value = 0.00309
```

```
#SARIMA(13,0,0)x(0,1,1)_12
```

```
model_26_train <- sarima(Avg_ExtentTS_Train, p=13, d=0, q=0, P=0, D=1, Q=1, S=12, details = TRUE)
```


```
## initial  value -0.794710  
## iter    2 value -1.008413  
## iter    3 value -1.149498  
## iter    4 value -1.247785  
## iter    5 value -1.312226  
## iter    6 value -1.340321  
## iter    7 value -1.362277  
## iter    8 value -1.368586  
## iter    9 value -1.372893  
## iter   10 value -1.376137  
## iter   11 value -1.377199  
## iter   12 value -1.377346  
## iter   13 value -1.377423  
## iter   14 value -1.377445  
## iter   15 value -1.377456  
## iter   16 value -1.377457  
## iter   17 value -1.377457  
## iter   18 value -1.377458  
## iter   19 value -1.377458
```

```
## iter 19 value -1.377458
## iter 19 value -1.377458
## final value -1.377458
## converged
## initial value -1.378066
## iter 2 value -1.378820
## iter 3 value -1.379100
## iter 4 value -1.379263
## iter 5 value -1.379282
## iter 6 value -1.379285
## iter 7 value -1.379285
## iter 8 value -1.379285
## iter 8 value -1.379285
## iter 8 value -1.379285
## final value -1.379285
## converged
```



```
model_26_train_residuals = resid(model_26_train$fit)
hist(model_26_train_residuals)
```

Histogram of model_26_train_residuals



```
shapiro.test(model_26_train_residuals)
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_26_train_residuals  
## W = 0.98804, p-value = 0.002459
```


```
#SARIMA(13,0,0)x(1,1,0)_12
```

```
model_27_train <- sarima(Avg_ExtentTS_Train, p=13, d=0, q=0, P=1, D=1, Q=0, S=12, details = TRUE)
```

```
## initial value -0.801164  
## iter 2 value -1.017847  
## iter 3 value -1.143844  
## iter 4 value -1.232274  
## iter 5 value -1.284447  
## iter 6 value -1.311521  
## iter 7 value -1.320540  
## iter 8 value -1.328806  
## iter 9 value -1.334262  
## iter 10 value -1.338005  
## iter 11 value -1.339163  
## iter 12 value -1.339774  
## iter 13 value -1.340389  
## iter 14 value -1.340593  
## iter 15 value -1.340695  
## iter 16 value -1.340719  
## iter 17 value -1.340730  
## iter 18 value -1.340734  
## iter 19 value -1.340734
```




```
## iter 20 value -1.340735
## iter 20 value -1.340735
## iter 20 value -1.340735
## final value -1.340735
## converged
## initial value -1.338080
## iter 2 value -1.338274
## iter 3 value -1.338681
## iter 4 value -1.338714
## iter 5 value -1.338795
## iter 6 value -1.338818
## iter 7 value -1.338842
## iter 8 value -1.338848
## iter 9 value -1.338850
## iter 10 value -1.338850
## iter 11 value -1.338850
## iter 11 value -1.338850
## iter 11 value -1.338850
## final value -1.338850
## converged
```



```
model_27_train_residuals = resid(model_27_train$fit)
hist(model_27_train_residuals)
```

Histogram of model_27_train_residuals



```
shapiro.test(model_27_train_residuals)
```


```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_27_train_residuals  
## W = 0.98468, p-value = 0.0003358
```

```
#SARIMA(13,0,0)x(0,1,3)_12
```

```
model_28_train <- sarima(Avg_ExtentTS_Train, p=13, d=0, q=0, P=0, D=1, Q=3, S=12 , details = TRUE)
```


```
## initial value -0.794710  
## iter 2 value -1.013936  
## iter 3 value -1.173142  
## iter 4 value -1.265645  
## iter 5 value -1.324556  
## iter 6 value -1.329399  
## iter 7 value -1.354647  
## iter 8 value -1.368673  
## iter 9 value -1.381611  
## iter 10 value -1.389125  
## iter 11 value -1.391039  
## iter 12 value -1.391755  
## iter 13 value -1.392861  
## iter 14 value -1.394528  
## iter 15 value -1.395490  
## iter 16 value -1.396737  
## iter 17 value -1.396990  
## iter 18 value -1.397072  
## iter 19 value -1.397085
```

```
## iter 20 value -1.397089
## iter 21 value -1.397089
## iter 22 value -1.397089
## iter 22 value -1.397089
## iter 22 value -1.397089
## final value -1.397089
## converged
## initial value -1.392610
## iter 2 value -1.392693
## iter 3 value -1.392839
## iter 4 value -1.392864
## iter 5 value -1.392891
## iter 6 value -1.392893
## iter 7 value -1.392897
## iter 8 value -1.392898
## iter 9 value -1.392903
## iter 10 value -1.392904
## iter 11 value -1.392908
## iter 12 value -1.392910
## iter 13 value -1.392911
## iter 14 value -1.392912
## iter 15 value -1.392912
## iter 16 value -1.392912
## iter 16 value -1.392912
## iter 16 value -1.392912
## final value -1.392912
## converged
```



```
model_28_train_residuals = resid(model_28_train$fit)
hist(model_28_train_residuals)
```

Histogram of model_28_train_residuals



```


shapiro.test(model_28_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_28_train_residuals
## W = 0.9863, p-value = 0.0008591
#SARIMA(13,0,0)x(0,1,4)_12
model_29_train <- sarima(Avg_ExtentTS_Train, p=13, d=0, q=0, P=0, D=1, Q=4, S=12 , details = TRUE)

## initial  value -0.794710
## iter    2 value -1.018370
## iter    3 value -1.185015
## iter    4 value -1.269162
## iter    5 value -1.338386
## iter    6 value -1.361375
## iter    7 value -1.373123
## iter    8 value -1.389868
## iter    9 value -1.390417
## iter   10 value -1.395151
## iter   11 value -1.396699
## iter   12 value -1.398587
## iter   13 value -1.399541
## iter   14 value -1.401027
## iter   15 value -1.401975
## iter   16 value -1.402455
## iter   17 value -1.402620
## iter   18 value -1.402660
## iter   19 value -1.402671
## iter   20 value -1.402672
## iter   21 value -1.402673
## iter   22 value -1.402673
## iter   22 value -1.402673
## iter   22 value -1.402673
## final  value -1.402673
## converged
## initial  value -1.397598
## iter    2 value -1.397676
## iter    3 value -1.397908
## iter    4 value -1.397918
## iter    5 value -1.397948
## iter    6 value -1.397953
## iter    7 value -1.397964
## iter    8 value -1.397968
## iter    9 value -1.397974
## iter   10 value -1.397981
## iter   11 value -1.397982
## iter   12 value -1.397982
## iter   13 value -1.397983
## iter   14 value -1.397983
## iter   14 value -1.397983
## iter   14 value -1.397983
## final  value -1.397983


```

converged



```
model_29_train_residuals = resid(model_29_train$fit)
hist(model_29_train_residuals)
```

Histogram of model_29_train_residuals



```
shapiro.test(model_29_train_residuals)
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_29_train_residuals  
## W = 0.98626, p-value = 0.0008373
```

```
#SARIMA(13,0,0)x(1,1,1)_12
```


```
model_30_train <- sarima(Avg_ExtentTS_Train, p=13, d=0, q=0, P=1, D=1, Q=1, S=12, details = TRUE)
```

```
## initial value -0.801164  
## iter 2 value -1.035924  
## iter 3 value -1.163984  
## iter 4 value -1.273600  
## iter 5 value -1.313108  
## iter 6 value -1.345098  
## iter 7 value -1.350268  
## iter 8 value -1.361690  
## iter 9 value -1.364074  
## iter 10 value -1.364704  
## iter 11 value -1.365211  
## iter 12 value -1.365454  
## iter 13 value -1.365788  
## iter 14 value -1.366112  
## iter 15 value -1.366710  
## iter 16 value -1.367095  
## iter 17 value -1.367549  
## iter 18 value -1.367600  
## iter 19 value -1.367614
```

```


## iter 20 value -1.367619
## iter 21 value -1.367620
## iter 22 value -1.367621
## iter 23 value -1.367621
## iter 23 value -1.367621
## iter 23 value -1.367621
## final value -1.367621
## converged
## initial value -1.378413
## iter 2 value -1.381305
## iter 3 value -1.381612
## iter 4 value -1.383730
## iter 5 value -1.384338
## iter 6 value -1.385117
## iter 7 value -1.385893
## iter 8 value -1.386815
## iter 9 value -1.387004
## iter 10 value -1.387072
## iter 11 value -1.387079
## iter 12 value -1.387081
## iter 13 value -1.387081
## iter 14 value -1.387081
## iter 14 value -1.387081
## iter 14 value -1.387081
## final value -1.387081
## converged

```




```
model_30_train_residuals = resid(model_30_train$fit)
hist(model_30_train_residuals)
```

Histogram of model_30_train_residuals



```
shapiro.test(model_30_train_residuals)
```


```
##
##  Shapiro-Wilk normality test
##
## data:  model_30_train_residuals
## W = 0.98623, p-value = 0.00082
```

```
#SARIMA(1,0,2)x(0,1,1)_12
```

```
model_31_train <- sarima(Avg_ExtentTS_Train, p=1, d=0, q=2, P=0, D=1, Q=1, S=12 , details = TRUE)
```


```
## initial  value -0.801086
## iter    2 value -1.138952
## iter    3 value -1.342237
## iter    4 value -1.355399
## iter    5 value -1.360693
## iter    6 value -1.362722
## iter    7 value -1.364221
## iter    8 value -1.365966
## iter    9 value -1.367305
## iter   10 value -1.367779
## iter   11 value -1.367834
## iter   12 value -1.367835
## iter   12 value -1.367835
## final   value -1.367835
## converged
## initial  value -1.365277
```

```
## iter 2 value -1.365339
## iter 3 value -1.365342
## iter 4 value -1.365342
## iter 5 value -1.365343
## iter 6 value -1.365344
## iter 7 value -1.365344
## iter 7 value -1.365344
## iter 7 value -1.365344
## final value -1.365344
## converged
```



```
model_31_train_residuals = resid(model_31_train$fit)
hist(model_31_train_residuals)
```

Histogram of model_31_train_residuals



```
shapiro.test(model_31_train_residuals)
```


```
##  
##  Shapiro-Wilk normality test  
##  
## data:  model_31_train_residuals  
## W = 0.98783, p-value = 0.002153
```

```
#SARIMA(2,0,1)x(0,1,1)_12
```

```
model_32_train <- sarima(Avg_ExtentTS_Train, p=2, d=0, q=1, P=0, D=1, Q=1, S=12, details = TRUE)
```


```
## initial  value -0.799783  
## iter    2 value -1.147624  
## iter    3 value -1.332335  
## iter    4 value -1.355514  
## iter    5 value -1.366626  
## iter    6 value -1.366675  
## iter    7 value -1.366841  
## iter    8 value -1.366857  
## iter    9 value -1.366899  
## iter   10 value -1.366998  
## iter   11 value -1.367081  
## iter   12 value -1.367109  
## iter   13 value -1.367111  
## iter   14 value -1.367112  
## iter   14 value -1.367112  
## iter   14 value -1.367112  
## final    value -1.367112  
## converged  
## initial  value -1.365359
```

```
## iter 2 value -1.365381
## iter 3 value -1.365409
## iter 4 value -1.365409
## iter 5 value -1.365410
## iter 6 value -1.365412
## iter 7 value -1.365416
## iter 8 value -1.365420
## iter 9 value -1.365421
## iter 10 value -1.365421
## iter 10 value -1.365421
## iter 10 value -1.365421
## final value -1.365421
## converged
```



```
model_32_train_residuals = resid(model_32_train$fit)
hist(model_32_train_residuals)
```

Histogram of model_32_train_residuals



```
shapiro.test(model_32_train_residuals)
```


```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_32_train_residuals  
## W = 0.98755, p-value = 0.001822
```

```
#SARIMA(2,0,2)x(0,1,1)_12
```

```
model_33_train <- sarima(Avg_ExtentTS_Train, p=2, d=0, q=2, P=0, D=1, Q=1, S=12 , details = TRUE)
```


```
## initial  value -0.799783  
## iter    2 value -1.066061  
## iter    3 value -1.228003  
## iter    4 value -1.328895  
## iter    5 value -1.356060  
## iter    6 value -1.366108  
## iter    7 value -1.366516  
## iter    8 value -1.367171  
## iter    9 value -1.367298  
## iter   10 value -1.367313  
## iter   11 value -1.367313  
## iter   11 value -1.367313  
## iter   11 value -1.367313  
## final    value -1.367313  
## converged  
## initial  value -1.365536  
## iter    2 value -1.365571  
## iter    3 value -1.365588  
## iter    4 value -1.365590
```

```
## iter 5 value -1.365591
## iter 6 value -1.365591
## iter 6 value -1.365591
## iter 6 value -1.365591
## final value -1.365591
## converged
```



```
model_33_train_residuals = resid(model_33_train$fit)
hist(model_33_train_residuals)
```

Histogram of model_33_train_residuals




```
shapiro.test(model_33_train_residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  model_33_train_residuals
## W = 0.98761, p-value = 0.001885

# Weird combo of optimal parameters for d=0, and d=1 (because had good results with d=D=1)
#SARIMA(1,0,1)x(3,1,0)_12
model_34_train <- sarima(Avg_ExtentTS_Train, p=1, d=0, q=1, P=3, D=1, Q=0, S=12, details = TRUE)


## initial  value -0.795069
## iter    2 value -1.107339
## iter    3 value -1.327476
## iter    4 value -1.351082
## iter    5 value -1.353717
## iter    6 value -1.360317
## iter    7 value -1.360498
## iter    8 value -1.360526
## iter    9 value -1.360547
## iter   10 value -1.360548
## iter   11 value -1.360548
## iter   11 value -1.360548
## iter   11 value -1.360548
## final   value -1.360548
## converged
## initial  value -1.361106
## iter    2 value -1.361247
## iter    3 value -1.361301
```

```
## iter 4 value -1.361318
## iter 5 value -1.361322
## iter 6 value -1.361323
## iter 7 value -1.361324
## iter 8 value -1.361324
## iter 8 value -1.361324
## iter 8 value -1.361324
## final value -1.361324
## converged
```



```
model_34_train_residuals = resid(model_34_train$fit)
hist(model_34_train_residuals)
```


Histogram of model_34_train_residuals




```
shapiro.test(model_34_train_residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  model_34_train_residuals
## W = 0.99035, p-value = 0.01064

# Weird combo of optimal parameters for d=0, and d=1 (because had good results with d=D=1)
#SARIMA(1,0,1)x(3,1,0)_12
model_35_train <- sarima(Avg_ExtentTS_Train, p=1, d=0, q=1, P=4, D=1, Q=0, S=12, details = TRUE)


## initial value -0.804892
## iter 2 value -1.078475
## iter 3 value -1.326207
## iter 4 value -1.348467
## iter 5 value -1.357442
## iter 6 value -1.366662
## iter 7 value -1.367642
## iter 8 value -1.367689
## iter 9 value -1.367707
## iter 10 value -1.367713
## iter 11 value -1.367715
## iter 12 value -1.367715
## iter 12 value -1.367715
## iter 12 value -1.367715
## final value -1.367715
## converged
## initial value -1.365081
## iter 2 value -1.365423
```

```
## iter 3 value -1.365564
## iter 4 value -1.365571
## iter 5 value -1.365575
## iter 6 value -1.365577
## iter 7 value -1.365577
## iter 8 value -1.365577
## iter 8 value -1.365577
## iter 8 value -1.365577
## final value -1.365577
## converged
```



```
model_35_train_residuals = resid(model_34_train$fit)
hist(model_35_train_residuals)
```

Histogram of model_35_train_residuals



```
shapiro.test(model_35_train_residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  model_35_train_residuals
## W = 0.99035, p-value = 0.01064
```

Summarize the fit of these models in a table.

```
library(huxtable)
goodness_of_fit <- hux(
  Model = c('SARIMA(1,0,1)x(0,1,1)_12', 'SARIMA(1,0,1)x(1,1,0)_12', 'SARIMA(1,0,1)x(0,1,3)_12', 'SARIMA(1,0,1)x(1,1,3)_12',
            'SARIMA(13,0,0)x(0,1,1)_12', 'SARIMA(13,0,0)x(1,1,0)_12', 'SARIMA(13,0,0)x(0,1,3)_12', 'SARIMA(13,0,0)x(1,1,3)_12',
            'SARIMA(1,0,2)x(0,1,1)_12', 'SARIMA(2,0,1)x(0,1,1)_12', 'SARIMA(2,0,2)x(0,1,1)_12', 'SARIMA(2,0,2)x(1,1,1)_12'),
  AIC = c(model_21_train$AIC, model_22_train$AIC, model_23_train$AIC, model_24_train$AIC, model_25_train$AIC, model_26_train$AIC, model_27_train$AIC, model_28_train$AIC, model_29_train$AIC, model_30_train$AIC, model_31_train$AIC, model_32_train$AIC, model_33_train$AIC, model_34_train$AIC, model_35_train$AIC),
  AICc = c(model_21_train$AICc, model_22_train$AICc, model_23_train$AICc, model_24_train$AICc, model_25_train$AICc, model_26_train$AICc, model_27_train$AICc, model_28_train$AICc, model_29_train$AICc, model_30_train$AICc, model_31_train$AICc, model_32_train$AICc, model_33_train$AICc, model_34_train$AICc, model_35_train$AICc),
  BIC = c(model_21_train$BIC, model_22_train$BIC, model_23_train$BIC, model_24_train$BIC, model_25_train$BIC, model_26_train$BIC, model_27_train$BIC, model_28_train$BIC, model_29_train$BIC, model_30_train$BIC, model_31_train$BIC, model_32_train$BIC, model_33_train$BIC, model_34_train$BIC, model_35_train$BIC),
  MSE = c(mean(model_21_train$residuals^2), mean(model_22_train$residuals^2), mean(model_23_train$residuals^2), mean(model_24_train$residuals^2), mean(model_25_train$residuals^2), mean(model_26_train$residuals^2), mean(model_27_train$residuals^2), mean(model_28_train$residuals^2), mean(model_29_train$residuals^2), mean(model_30_train$residuals^2), mean(model_31_train$residuals^2), mean(model_32_train$residuals^2), mean(model_33_train$residuals^2), mean(model_34_train$residuals^2), mean(model_35_train$residuals^2))
)

goodness_of_fit %>%
  set_number_format(col=c(2,3,4,5), value=3) %>%
```

```

set_bottom_border(c(1,11,14), everywhere) %>%
set_bold(c(2,4,5,6), everywhere) %>%
set_background_color(evens, everywhere, "grey95")


```

Model	AIC	AICc	BIC	MSE
SARIMA(1,0,1)x(0,1,1)_12	0.133	0.134	0.185	0.061
SARIMA(1,0,1)x(1,1,0)_12	0.250	0.250	0.301	0.070
SARIMA(1,0,1)x(0,1,3)_12	0.137	0.137	0.209	0.061
SARIMA(1,0,1)x(0,1,4)_12	0.134	0.135	0.216	0.060
SARIMA(1,0,1)x(1,1,1)_12	0.138	0.138	0.200	0.061
SARIMA(13,0,0)x(0,1,1)_12	0.163	0.166	0.327	0.060
SARIMA(13,0,0)x(1,1,0)_12	0.244	0.247	0.408	0.066
SARIMA(13,0,0)x(0,1,3)_12	0.146	0.150	0.331	0.058
SARIMA(13,0,0)x(0,1,4)_12	0.141	0.146	0.336	0.057
SARIMA(13,0,0)x(1,1,1)_12	0.152	0.156	0.327	0.059
SARIMA(1,0,2)x(0,1,1)_12	0.138	0.139	0.200	0.061
SARIMA(2,0,1)x(0,1,1)_12	0.138	0.139	0.200	0.061
SARIMA(2,0,2)x(0,1,1)_12	0.143	0.144	0.215	0.061
SARIMA(1,0,1)x(3,1,0)_12	0.152	0.152	0.224	0.062
SARIMA(1,0,1)x(4,1,0)_12	0.148	0.149	0.231	0.062


Model Selection

We evaluate performance on the test set of a few of the models which gave the best fit.


```
model_21_train_forecast<- sarima.for(Avg_ExtentTS_Train, n.ahead=22, p=1,d=0,q=1,P=0,D=1,Q=1,S=12)  
lines(Avg_ExtentTS_Test,col='blue',type='b',pch=16)
```




```
model_23_train_forecast<- sarima.for(Avg_ExtentTS_Train, n.ahead=22, p=1,d=0,q=1,P=0,D=1,Q=3,S=12)  
lines(Avg_ExtentTS_Test,col='blue',type='b',pch=16)
```



```
model_24_train_forecast<- sarima.for(Avg_ExtentTS_Train, n.ahead=22, p=1,d=0,q=1,P=0,D=1,Q=4,S=12)
lines(Avg_ExtentTS_Test,col='blue',type='b',pch=16)
```



```
model_25_train_forecast<- sarima.for(Avg_ExtentTS_Train, n.ahead=22, p=1,d=0,q=1,P=1,D=1,Q=1,S=12)
lines(Avg_ExtentTS_Test,col='blue',type='b',pch=16)
```



```
mean((model_21_train_forecast$pred-Avg_ExtentTS_Test)^2)
```

```
## [1] 0.3343123
```

```
mean((model_23_train_forecast$pred-Avg_ExtentTS_Test)^2)
```

```
## [1] 0.3455438
```

```
mean((model_24_train_forecast$pred-Avg_ExtentTS_Test)^2)
```

```
## [1] 0.3370083
```

```
mean((model_25_train_forecast$pred-Avg_ExtentTS_Test)^2)
```

```
## [1] 0.3340769
```

Summarize these results in a table.

```
sarima_prediction <- hux(
  Model = c('SARIMA(1,0,1)x(0,1,1)_12', 'SARIMA(1,0,1)x(0,1,3)_12', 'SARIMA(1,0,1)x(0,1,4)_12', 'SARIMA(1,0,1)x(0,1,5)_12'),
  PMSE = c(mean((model_21_train_forecast$pred-Avg_ExtentTS_Test)^2), mean((model_23_train_forecast$pred-Avg_ExtentTS_Test)^2),
    mean((model_24_train_forecast$pred-Avg_ExtentTS_Test)^2), mean((model_25_train_forecast$pred-Avg_ExtentTS_Test)^2))


sarima_prediction %>%
  set_number_format(col=2, value=3) %>%
  set_bottom_border(1, everywhere) %>%
  set_background_color(evens, everywhere, "grey95")
```


Model	PMSE
SARIMA(1,0,1)x(0,1,1)_12	0.334
SARIMA(1,0,1)x(0,1,3)_12	0.346
SARIMA(1,0,1)x(0,1,4)_12	0.337
SARIMA(1,0,1)x(1,1,1)_12	0.334

```
# Best SARIMA model
```

```
model_21 <- sarima(Avg_ExtentTS, p=1, d=0, q=1, P=0, D=1, Q=1, S=12 , details = TRUE)
```


```
## initial value -0.789422
## iter 2 value -1.114422
## iter 3 value -1.355131
## iter 4 value -1.356779
## iter 5 value -1.365174
## iter 6 value -1.367226
## iter 7 value -1.367725
## iter 8 value -1.367797
## iter 9 value -1.367800
## iter 10 value -1.367803
## iter 11 value -1.367803
## iter 12 value -1.367804
## iter 12 value -1.367804
## iter 12 value -1.367804
## final value -1.367804
## converged
## initial value -1.365801
## iter 2 value -1.365834
## iter 3 value -1.365859
## iter 4 value -1.365875
## iter 5 value -1.365876
## iter 5 value -1.365876
## iter 5 value -1.365876
## final value -1.365876
## converged
```



```
model_21_residuals = resid(model_21$fit)
shapiro.test(model_21_residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  model_21_residuals
## W = 0.9883, p-value = 0.001968
```


```
# Getting predictions for best SARIMA and plotting
pred_model_21 = sarima.for(Avg_ExtentTS, n.ahead=60, p=1, d=0, q=1, P=0, D=1, Q=1, S=12)
```



```
Upper_Limit = pred_model_21$pred + 2*pred_model_21$se # upper prediction band
Lower_Limit = pred_model_21$pred - 2*pred_model_21$se # lower prediction band
plot(Avg_ExtentTS , xlim = c(2015 , 2027+10/12), ylab="Average Ice Extent", main="Prediction from SARIM

#The three lines below plot the prediction interval in a grey scale
x = c(time(Upper_Limit) , rev(time(Upper_Limit)))
y = c(Upper_Limit , rev(Lower_Limit))
polygon(x, y, col="grey", border=NA)

#The three line below add the predicted values and highlight the borders of the prediction interval
lines(Upper_Limit, col="black" , lty=2)
lines(Lower_Limit, col="black", lty=2)
lines(pred_model_21$pred , col="red")
abline(v=2022+9.5/12, lty="dashed")
```



```
model_21_residuals = resid(model_21$fit)
shapiro.test(model_21_residuals)
```

```
##
## Shapiro-Wilk normality test
##
## data:  model_21_residuals
## W = 0.9883, p-value = 0.001968
```

Summary of prediction from best sarima model.

```
summary_prediction <- summary(pred_model_21$pred)
summary <- as.data.frame(rbind(matrix(summary_prediction,nrow=1)))
colnames(summary) <- c("Min", "First Quartile", "Median", "Mean", "Third Quartile", "Max")

summary_table <- hux(summary)

summary_table %>%
  set_number_format(3) %>%
  set_align(everywhere, everywhere, "center") %>%
  set_bottom_border(1, everywhere) %>%
  set_bold(1, everywhere)
```

Min	First Quartile	Median	Mean	Third Quartile	Max
4.257	7.148	11.065	10.197	13.431	14.442

First year of prediction from best sarima model.

```
prediction <- as.data.frame(rbind(matrix(head(pred_model_21$pred, 12),nrow=1)))
colnames(prediction) <- c("Nov.", "Dec.", "Jan.", "Feb.", "Mar.", "Apr.", "May", "Jun.", "Jul.", "Aug.",
                           "Sep.", "Oct.")

prediction_table <- hux(prediction)


prediction_table %>%
  set_number_format(3) %>%
  set_align(everywhere, everywhere, "center") %>%
  set_bottom_border(1, everywhere) %>%
  set_bold(1, everywhere) %>%
  print_latex(tabular_only=TRUE)

##
##
## ```{=latex}
##
## \providecommand{\huxb}[2]{\arrayrulecolor[RGB]{#1}\global\arrayrulewidth=#2pt}
## \providecommand{\huxvb}[2]{\color[RGB]{#1}\vrule width #2pt}
## \providecommand{\huxtpad}[1]{\rule{0pt}{#1}}
## \providecommand{\huxbpad}[1]{\rule[-#1]{0pt}{#1}}
## \begin{tabular}{l l l l l l l l l l l}
##
##
## \hhline{}
## \arrayrulecolor{black}
##
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}c!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Dec.} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Jan.} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Feb.} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Mar.} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Apr.} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{May} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Jun.} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Jul.} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Aug.} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Sep.} \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} \textbf{Oct.} \hspace{6pt}}
##
##
## \hhline{>{\huxb{0, 0, 0}{0.4}}->{\huxb{0, 0, 0}{0.4}}->{\huxb{0, 0, 0}{0.4}}->{\huxb{0, 0, 0}{0.4}}->}
## \arrayrulecolor{black}
##
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}c!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 11.962 \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 13.526 \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 14.323 \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 14.442 \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 13.656 \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 12.376 \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 10.562 \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 7.691 \hspace{6pt}}
## \multicolumn{1}{!{\huxvb{0, 0, 0}{0}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 5.381 \hspace{6pt}}
```


```
## \multicolumn{1}{c!{\huxvb{0, 0, 0}{0}}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 4.468 \hspace{6pt}
## \multicolumn{1}{c!{\huxvb{0, 0, 0}{0}}}{\huxtpad{6pt + 1em}\centering \hspace{6pt} 6.168 \hspace{6pt}
##
##
## \hhline{}
## \arrayrulecolor{black}
## \end{tabular}
## ``
```

```
# code used for additional plots in report
```

```
plot(acf(Avg_ExtentTS, lag.max=36), main="ACF of Aggregated Data")
```




ACF of Aggregated Data




```
plot(acf(mlr_train_3$residuals, lag.max=36), main="ACF of MLR Residuals")
```

Series mlr_train_3\$residuals




ACF of MLR Residuals




```
plot(acf(residuals_HW, lag.max=36), main="ACF of HW Residuals")
```

Series residuals_HW




ACF of HW Residuals




```
plot(acf(diff12.diff.Extent, lag.max=36), main="ACF of Differenced Data (Seasonal+Regular)")
```

Series diff12.diff.Extent




ACF of Differenced Data (Seasonal+Regular)




```
plot(pacf(diff12.diff.Extent, lag.max=36), main="PACF of Differenced Data (Seasonal+Regular)")
```

Series diff12.diff.Extent




PACF of Differenced Data (Seasonal+Regular)




```
plot(acf(diff12.Extent, lag.max=36), main="ACF of Differenced Data (Seasonal)")
```

Series diff12.Extent




ACF of Differenced Data (Seasonal)




```
plot(pacf(diff12.Extent, lag.max=36), main="PACF of Differenced Data (Seasonal)")
```

Series diff12.Extent



PACF of Differenced Data (Seasonal)




Things we tried but didn't make the report

Unaggregated Data

We analyze the unaggregated data.


```
# make a TS object
ExtentTS <- ts(df$extent, frequency=365, start=year(df$YYMMDD[1]))

plot(ExtentTS)
```



```
acf(ExtentTS, lag.max=365)
```

Series ExtentTS



Variance

From the plot, we see a clear seasonal pattern, and perhaps a decreasing linear trend.

It is unclear whether variance is constant. We test this using the Fligner-Keileen test.

```
# do Fligner test for constant variance.
segments = factor(c(rep(1:4, each=2542), rep(5, times=2543)))
fligner.test(ExtentTS, segments)
```

```
##
## Fligner-Killeen test of homogeneity of variances
##
## data: ExtentTS and segments
## Fligner-Killeen:med chi-squared = 219.12, df = 9, p-value < 2.2e-16
segments = factor(c(rep(1:49, each=254), rep(50, times=265)))
fligner.test(ExtentTS, segments)
```

```
##
## Fligner-Killeen test of homogeneity of variances
##
## data: ExtentTS and segments
## Fligner-Killeen:med chi-squared = 1326.7, df = 49, p-value < 2.2e-16
segments = factor(c(rep(1:99, each=127), rep(100, times=138)))
fligner.test(ExtentTS, segments)
```

```
##
## Fligner-Killeen test of homogeneity of variances
##
## data: ExtentTS and segments
## Fligner-Killeen:med chi-squared = 3386.9, df = 99, p-value < 2.2e-16
segments = factor(c(rep(1:34, each=364), rep(35, times=335))) # corresp
fligner.test(ExtentTS, segments)
```


```
##
## Fligner-Killeen test of homogeneity of variances
##
## data: ExtentTS and segments
## Fligner-Killeen:med chi-squared = 368.79, df = 34, p-value < 2.2e-16
```

All give really low p-value so may conclude that variance is not constant. However, this could be due to the amount of data we have.

Try to remove non-stationarity using Regression (Multiple Linear, Ridge, Lasso, Elastic Net).

```
m1r <- lm(ExtentTS~time(ExtentTS)+factor(cycle(ExtentTS)))  
#summary(m1r) #verrrrryyyyyy long output and complicated model.
```

```
plot(ExtentTS)
points(time(ExtentTS),predict.lm(mlr),type='l',col='red')
```




We see from the regression, that including daily data leads to a very complicated regression model, and acf plot which has to go way beyond recommended lag to observe an entire period. For this reason, and because we care mostly about overall trend and not daily fluctuation, we proceed with the aggregated data.

Differencing on Entire Data


Try differencing to remove non-stationarity.

```
acf(Avg_ExtentTS, lag.max=36)
```


Series Avg_ExtentTS




```
plot(Avg_ExtentTS)
```




```
#differencing in lag of season  
diff12.Extent=diff(Avg_ExtentTS, lag=12)  
acf(diff12.Extent, lag.max=36)
```

Series diff12.Extent




```
pacf(diff12.Extent, lag.max=36)
```

Series diff12.Extent




```
plot(diff12.Extent)
```




```
#regular differencing
diff.Extent=diff(Avg_ExtentTS)
acf(diff.Extent, lag.max=36)
```

Series diff.Extent




```
plot(diff.Extent)
```




```
#seasonal+regular differencing
diff12.diff.Extent=diff(diff12.Extent)
acf(diff12.diff.Extent, lag.max=36)
```

Series diff12.diff.Extent




```
plot(diff12.diff.Extent)
```




```
acf(diff12.diff.Extent, lag.max=36)
```

Series diff12.diff.Extent



```
pacf(diff12.diff.Extent, lag.max=36)
```




Noted after that differencing should have been done on training data.

Smoothing, followed by differencing


Try smoothing before differencing to see effect on acf

```
smoothing <- HoltWinters(Avg_ExtentTS_Train, season="additive")
smoothed <- smoothing$fitted[,1]
diff12.Extent_smooth=diff(smoothed, lag=12)
acf(diff12.Extent_smooth, lag.max=36)
```

Series diff12.Extent_smooth




```
plot(diff12.Extent_smooth)
```




```
diff12.diff.Extent_smooth=diff(diff12.Extent_smooth)  
acf(diff12.diff.Extent_smooth, lag.max=36)
```

Series diff12.diff.Extent_smooth




```
plot(diff12.diff.Extent_smooth)
```




```
acf(diff12.diff.Extent_smooth, lag.max=36)
```


Series diff12.diff.Extent_smooth



```
pacf(diff12.diff.Extent_smooth, lag.max=36)
```

Series diff12.diff.Extent_smooth




This didn't change anything.

Model Fitting

Seasonal + regular differencing.


```
#SARIMA(0,1,2)x(0,1,1)_12
model_1_train <- sarima(Avg_ExtentTS_Train, p=0, d=1, q=2, P=0, D=1, Q=1, S=12 , details = TRUE)

## initial  value -1.046835
## iter    2 value -1.270895
## iter    3 value -1.286671
## iter    4 value -1.288786
## iter    5 value -1.299437
## iter    6 value -1.301999
## iter    7 value -1.302786
## iter    8 value -1.302845
## iter    9 value -1.302850
## iter   10 value -1.302854
## iter   11 value -1.302854
## iter   11 value -1.302854
## iter   11 value -1.302854
## final   value -1.302854
## converged
## initial  value -1.299104
## iter    2 value -1.299174
## iter    3 value -1.299223
## iter    4 value -1.299225
## iter    5 value -1.299225
## iter    5 value -1.299225
## iter    5 value -1.299225
## final   value -1.299225
## converged
```



```
model_1_train_residuals = resid(model_1_train$fit)
hist(model_1_train_residuals)
```

Histogram of model_1_train_residuals



```


shapiro.test(model_1_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_1_train_residuals
## W = 0.98032, p-value = 3.21e-05

#SARIMA(0,1,2)x(3,1,0)_12
model_2_train <- sarima(Avg_ExtentTS_Train, p=0, d=1, q=2, P=3, D=1, Q=0, S=12 , details = TRUE)


## initial  value -1.039490
## iter    2 value -1.176658
## iter    3 value -1.269525
## iter    4 value -1.294457
## iter    5 value -1.299214
## iter    6 value -1.299421
## iter    7 value -1.299429
## iter    8 value -1.299429
## iter    8 value -1.299429
## iter    8 value -1.299429
## final   value -1.299429
## converged
## initial  value -1.297378
## iter    2 value -1.297439
## iter    3 value -1.297448
## iter    4 value -1.297451
## iter    5 value -1.297451
## iter    6 value -1.297451
## iter    6 value -1.297451
## iter    6 value -1.297451
## final   value -1.297451
## converged

```



```
model_2_train_residuals = resid(model_2_train$fit)
hist(model_2_train_residuals)
```

Histogram of model_2_train_residuals



```


shapiro.test(model_2_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_2_train_residuals
## W = 0.98446, p-value = 0.0002963

#SARIMA(0,1,2)x(1,1,1)_12
model_3_train <- sarima(Avg_ExtentTS_Train, p=0, d=1, q=2, P=1, D=1, Q=1, S=12 , details = TRUE)


## initial  value -1.044499
## iter    2 value -1.237721
## iter    3 value -1.282334
## iter    4 value -1.287558
## iter    5 value -1.293725
## iter    6 value -1.293955
## iter    7 value -1.294047
## iter    8 value -1.294078
## iter    9 value -1.294083
## iter   10 value -1.294083
## iter   10 value -1.294083
## final   value -1.294083
## converged
## initial  value -1.299017
## iter    2 value -1.299579
## iter    3 value -1.300082
## iter    4 value -1.300197
## iter    5 value -1.300210
## iter    6 value -1.300211
## iter    7 value -1.300211
## iter    8 value -1.300211
## iter    8 value -1.300211
## iter    8 value -1.300211
## final   value -1.300211
## converged

```



```
model_3_train_residuals = resid(model_3_train$fit)
hist(model_3_train_residuals)
```

Histogram of model_3_train_residuals



```

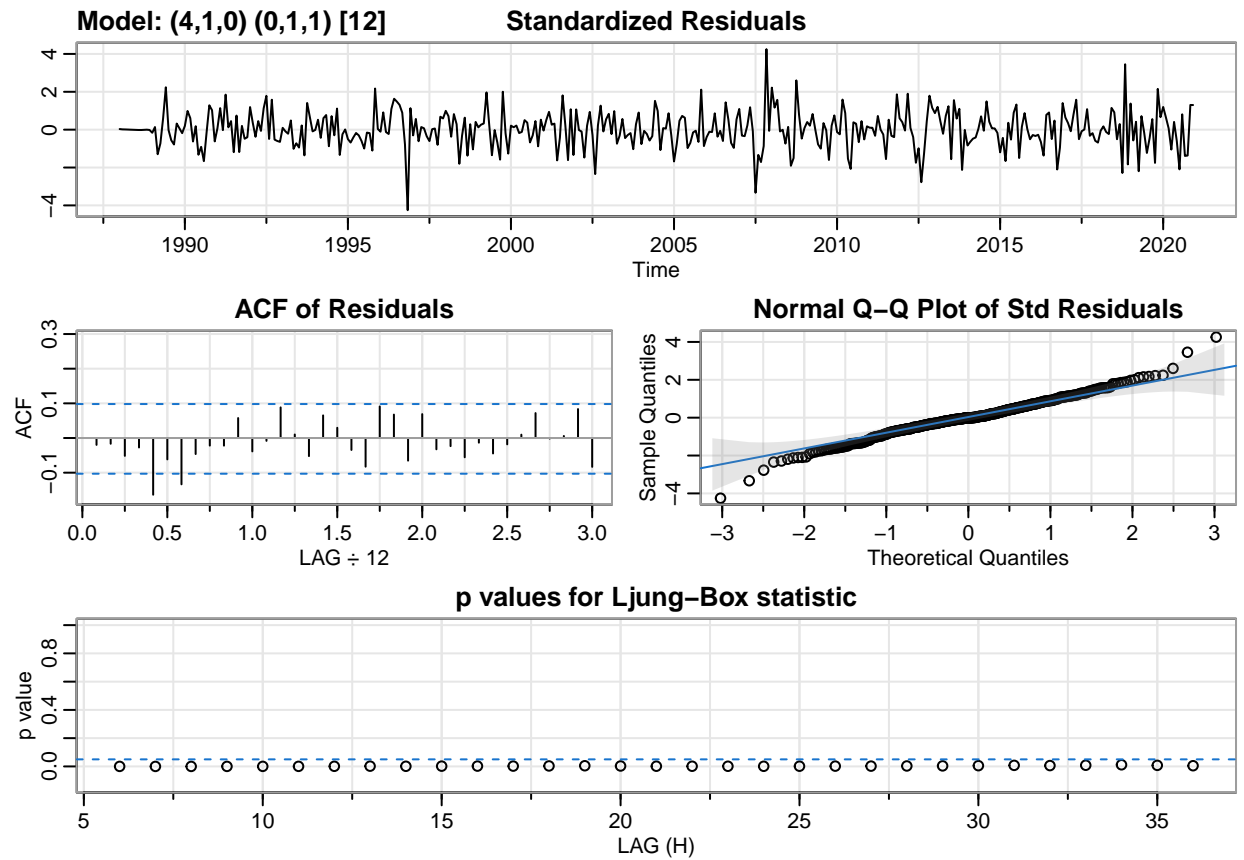
shapiro.test(model_3_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_3_train_residuals
## W = 0.98178, p-value = 6.856e-05

#SARIMA(4,1,0)x(0,1,1)_12
model_4_train <- sarima(Avg_ExtentTS_Train, p=4, d=1, q=0, P=0, D=1, Q=1, S=12 , details = TRUE)


## initial  value -1.045273
## iter    2 value -1.275376
## iter    3 value -1.289687
## iter    4 value -1.291649
## iter    5 value -1.300383
## iter    6 value -1.302661
## iter    7 value -1.303220
## iter    8 value -1.303284
## iter    9 value -1.303285
## iter    9 value -1.303285
## iter    9 value -1.303285
## final   value -1.303285
## converged
## initial  value -1.305545
## iter    2 value -1.305718
## iter    3 value -1.305759
## iter    4 value -1.305764
## iter    5 value -1.305764
## iter    5 value -1.305764
## iter    5 value -1.305764
## final   value -1.305764
## converged

```

```
model_4_train_residuals = resid(model_4_train$fit)
hist(model_4_train_residuals)
```

Histogram of model_4_train_residuals




```

shapiro.test(model_4_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_4_train_residuals
## W = 0.98293, p-value = 0.0001267
#SARIMA(4,1,0)x(3,1,0)_12
model_5_train <- sarima(Avg_ExtentTS_Train, p=4, d=1, q=0, P=3, D=1, Q=0, S=12 , details = TRUE)


## initial  value -1.034768
## iter    2 value -1.176278
## iter    3 value -1.267070
## iter    4 value -1.292224
## iter    5 value -1.298051
## iter    6 value -1.298311
## iter    7 value -1.298325
## iter    8 value -1.298325
## iter    8 value -1.298325
## iter    8 value -1.298325
## final   value -1.298325
## converged
## initial  value -1.301857
## iter    2 value -1.302141
## iter    3 value -1.302152
## iter    4 value -1.302153
## iter    5 value -1.302153
## iter    5 value -1.302153
## iter    5 value -1.302153
## final   value -1.302153
## converged

```



```
model_5_train_residuals = resid(model_5_train$fit)
hist(model_5_train_residuals)
```

Histogram of model_5_train_residuals



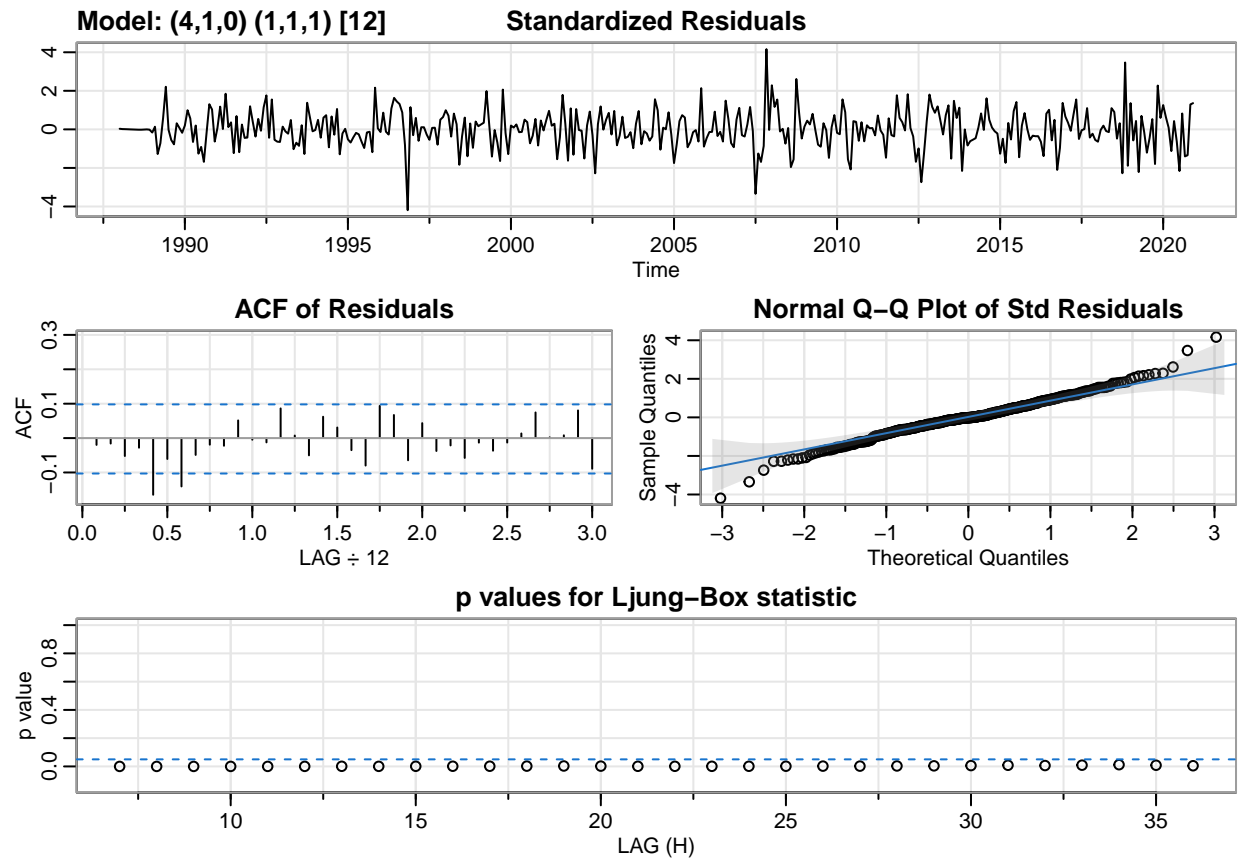
```

shapiro.test(model_5_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_5_train_residuals
## W = 0.98537, p-value = 0.0004979
#SARIMA(4,1,0)x(1,1,1)_12
model_6_train <- sarima(Avg_ExtentTS_Train, p=4, d=1, q=0, P=1, D=1, Q=1, S=12 , details = TRUE)


## initial  value -1.041584
## iter    2 value -1.241588
## iter    3 value -1.287381
## iter    4 value -1.292821
## iter    5 value -1.300492
## iter    6 value -1.300777
## iter    7 value -1.300797
## iter    8 value -1.300801
## iter    9 value -1.300802
## iter    9 value -1.300802
## iter    9 value -1.300802
## final   value -1.300802
## converged
## initial  value -1.305672
## iter    2 value -1.306365
## iter    3 value -1.306596
## iter    4 value -1.306681
## iter    5 value -1.306698
## iter    6 value -1.306701
## iter    7 value -1.306701
## iter    7 value -1.306701
## iter    7 value -1.306701
## final   value -1.306701
## converged

```



```
model_6_train_residuals = resid(model_6_train$fit)
hist(model_6_train_residuals)
```

Histogram of model_6_train_residuals




```

shapiro.test(model_6_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_6_train_residuals
## W = 0.98408, p-value = 0.0002389
#SARIMA(5,1,0)x(0,1,1)_12
model_7_train <- sarima(Avg_ExtentTS_Train, p=5, d=1, q=0, P=0, D=1, Q=1, S=12 , details = TRUE)


## initial  value -1.051141
## iter    2 value -1.280688
## iter    3 value -1.294984
## iter    4 value -1.299314
## iter    5 value -1.306323
## iter    6 value -1.308043
## iter    7 value -1.308352
## iter    8 value -1.308378
## iter    9 value -1.308378
## iter    9 value -1.308378
## iter    9 value -1.308378
## final   value -1.308378
## converged
## initial  value -1.312741
## iter    2 value -1.313183
## iter    3 value -1.313486
## iter    4 value -1.313497
## iter    5 value -1.313500
## iter    6 value -1.313500
## iter    6 value -1.313500
## iter    6 value -1.313500
## final   value -1.313500
## converged

```



```
model_7_train_residuals = resid(model_7_train$fit)
hist(model_7_train_residuals)
```

Histogram of model_7_train_residuals



```


shapiro.test(model_7_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_7_train_residuals
## W = 0.98416, p-value = 0.0002503

#SARIMA(5,1,0)x(3,1,0)_12
model_8_train <- sarima(Avg_ExtentTS_Train, p=5, d=1, q=0, P=3, D=1, Q=0, S=12 , details = TRUE)


## initial  value -1.034323
## iter    2 value -1.181171
## iter    3 value -1.276196
## iter    4 value -1.303175
## iter    5 value -1.309458
## iter    6 value -1.309745
## iter    7 value -1.309762
## iter    8 value -1.309762
## iter    8 value -1.309762
## iter    8 value -1.309762
## final   value -1.309762
## converged
## initial  value -1.309715
## iter    2 value -1.310249
## iter    3 value -1.310286
## iter    4 value -1.310300
## iter    5 value -1.310300
## iter    6 value -1.310300
## iter    6 value -1.310300
## iter    6 value -1.310300
## final   value -1.310300
## converged

```

```
model_8_train_residuals = resid(model_8_train$fit)
hist(model_8_train_residuals)
```

Histogram of model_8_train_residuals



```


shapiro.test(model_8_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_8_train_residuals
## W = 0.98765, p-value = 0.00193

#SARIMA(5,1,0)x(1,1,1)_12
model_9_train <- sarima(Avg_ExtentTS_Train, p=5, d=1, q=0, P=1, D=1, Q=1, S=12 , details = TRUE)


## initial  value -1.047019
## iter    2 value -1.250977
## iter    3 value -1.298091
## iter    4 value -1.305471
## iter    5 value -1.315008
## iter    6 value -1.315491
## iter    7 value -1.315527
## iter    8 value -1.315540
## iter    9 value -1.315541
## iter    9 value -1.315541
## iter    9 value -1.315541
## final   value -1.315541
## converged
## initial  value -1.313961
## iter    2 value -1.314344
## iter    3 value -1.314392
## iter    4 value -1.314401
## iter    5 value -1.314406
## iter    6 value -1.314406
## iter    7 value -1.314406
## iter    7 value -1.314406
## iter    7 value -1.314406
## final   value -1.314406
## converged

```



```
model_9_train_residuals = resid(model_9_train$fit)
hist(model_9_train_residuals)
```

Histogram of model_9_train_residuals



```
shapiro.test(model_9_train_residuals)
```


```
##  
##  Shapiro-Wilk normality test  
##  
## data:  model_9_train_residuals  
## W = 0.98544, p-value = 0.0005193
```

```
#SARIMA(1,1,1)x(0,1,1)_12
```

```
model_10_train <- sarima(Avg_ExtentTS_Train, p=1, d=1, q=1, P=0, D=1, Q=1, S=12 , details = TRUE)
```


```
## initial  value -1.045552  
## iter    2 value -1.242919  
## iter    3 value -1.258615  
## iter    4 value -1.265606  
## iter    5 value -1.269085  
## iter    6 value -1.271594  
## iter    7 value -1.272018  
## iter    8 value -1.272061  
## iter    9 value -1.272075  
## iter   10 value -1.272103  
## iter   11 value -1.273789  
## iter   12 value -1.274147  
## iter   13 value -1.274999  
## iter   14 value -1.275080  
## iter   15 value -1.276771  
## iter   16 value -1.276973  
## iter   17 value -1.277440  
## iter   18 value -1.277524  
## iter   19 value -1.277528  
## iter   20 value -1.277529  
## iter   21 value -1.277529  
## iter   22 value -1.277530  
## iter   23 value -1.277530  
## iter   24 value -1.277530  
## iter   25 value -1.277531  
## iter   26 value -1.277531  
## iter   26 value -1.277531  
## iter   26 value -1.277531  
## final   value -1.277531  
## converged  
## initial  value -1.276765  
## iter    2 value -1.276787  
## iter    3 value -1.276795  
## iter    4 value -1.276795  
## iter    5 value -1.276795  
## iter    6 value -1.276796  
## iter    7 value -1.276797  
## iter    8 value -1.276797  
## iter    9 value -1.276798  
## iter   10 value -1.276798  
## iter   10 value -1.276798  
## iter   10 value -1.276798  
## final   value -1.276798
```

converged



```
model_10_train_residuals = resid(model_10_train$fit)
hist(model_10_train_residuals)
```

Histogram of model_10_train_residuals



```
shapiro.test(model_10_train_residuals)
```

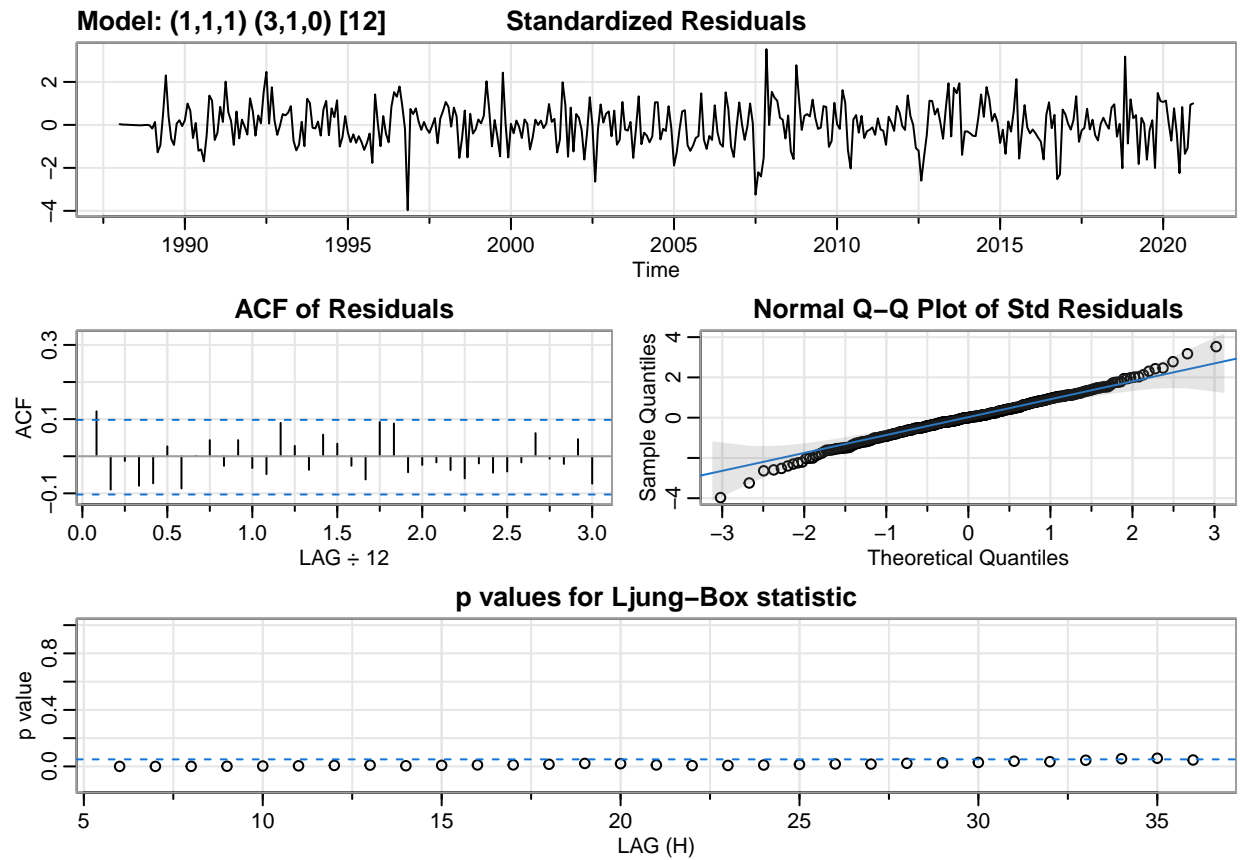
```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_10_train_residuals  
## W = 0.97359, p-value = 1.319e-06
```

```
#SARIMA(1,1,1)x(3,1,0)_12
```

```
model_11_train <- sarima(Avg_ExtentTS_Train, p=1, d=1, q=1, P=3, D=1, Q=0, S=12, details = TRUE)
```


```
## initial value -1.038097  
## iter 2 value -1.144641  
## iter 3 value -1.239216  
## iter 4 value -1.265137  
## iter 5 value -1.269506  
## iter 6 value -1.269577  
## iter 7 value -1.269587  
## iter 8 value -1.269588  
## iter 9 value -1.269590  
## iter 10 value -1.269601  
## iter 11 value -1.269726  
## iter 12 value -1.269928  
## iter 13 value -1.269990  
## iter 14 value -1.270074  
## iter 15 value -1.270123  
## iter 16 value -1.270398  
## iter 17 value -1.274020  
## iter 18 value -1.276115  
## iter 19 value -1.281384
```

```
## iter 20 value -1.292576
## iter 21 value -1.304600
## iter 22 value -1.318963
## iter 23 value -1.331206
## iter 24 value -1.334219
## iter 25 value -1.334457
## iter 26 value -1.334997
## iter 27 value -1.336065
## iter 28 value -1.336081
## iter 29 value -1.336093
## iter 30 value -1.336095
## iter 30 value -1.336095
## iter 30 value -1.336095
## final value -1.336095
## converged
## initial value -1.332968
## iter 2 value -1.334981
## iter 3 value -1.336741
## iter 4 value -1.336979
## iter 5 value -1.337236
## iter 6 value -1.337267
## iter 7 value -1.337272
## iter 8 value -1.337273
## iter 9 value -1.337273
## iter 9 value -1.337273
## iter 9 value -1.337273
## final value -1.337273
## converged
```



```
model_11_train_residuals = resid(model_11_train$fit)
hist(model_11_train_residuals)
```

Histogram of model_11_train_residuals




```

shapiro.test(model_11_train_residuals)


##
##  Shapiro-Wilk normality test
##
## data:  model_11_train_residuals
## W = 0.99044, p-value = 0.01131

#SARIMA(1,1,1)x(1,1,1)_12
model_12_train <- sarima(Avg_ExtentTS_Train, p=1, d=1, q=1, P=1, D=1, Q=1, S=12 , details = TRUE)

## initial  value -1.043857
## iter    2 value -1.216764
## iter    3 value -1.259588
## iter    4 value -1.264738
## iter    5 value -1.271016
## iter    6 value -1.271155
## iter    7 value -1.271164
## iter    8 value -1.271165
## iter    9 value -1.271166
## iter   10 value -1.271174
## iter   11 value -1.271177
## iter   12 value -1.271179
## iter   13 value -1.271182
## iter   14 value -1.271192
## iter   15 value -1.271226
## iter   16 value -1.271350
## iter   17 value -1.271513
## iter   18 value -1.271902
## iter   19 value -1.272301
## iter   20 value -1.272695
## iter   21 value -1.272800
## iter   22 value -1.272856
## iter   23 value -1.272922
## iter   24 value -1.273079
## iter   25 value -1.273307
## iter   26 value -1.273488
## iter   27 value -1.273600
## iter   28 value -1.273602
## iter   29 value -1.273604
## iter   30 value -1.273604
## iter   31 value -1.273604
## iter   31 value -1.273604
## iter   31 value -1.273604
## final   value -1.273604
## converged
## initial  value -1.276218
## iter    2 value -1.276807
## iter    3 value -1.277527
## iter    4 value -1.277549
## iter    5 value -1.277577
## iter    6 value -1.277709
## iter    7 value -1.277802
## iter    8 value -1.277869


```

```
## iter 9 value -1.277873
## iter 10 value -1.277874
## iter 10 value -1.277874
## final value -1.277874
## converged
```



```
model_12_train_residuals = resid(model_4_train$fit)
hist(model_12_train_residuals)
```

Histogram of model_12_train_residuals



```
shapiro.test(model_12_train_residuals)
```

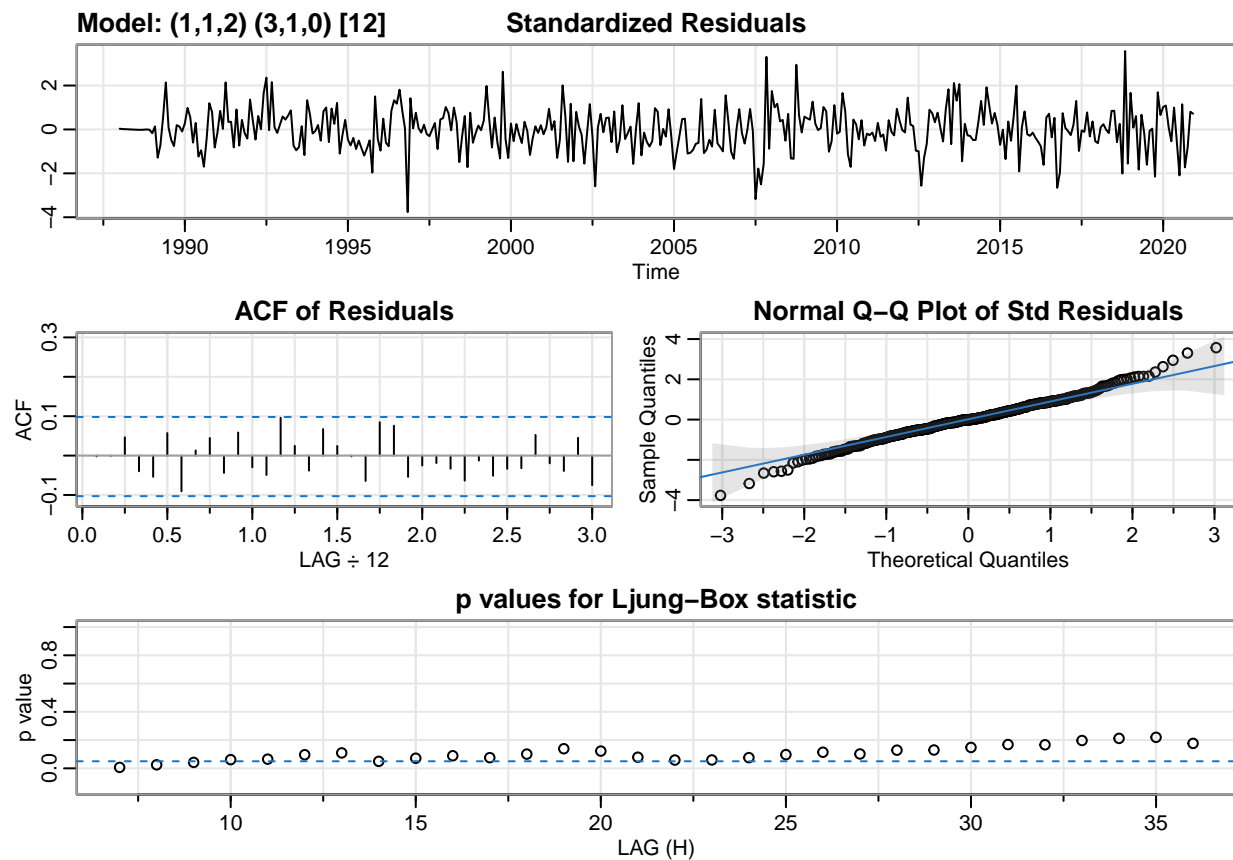
```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_12_train_residuals  
## W = 0.98293, p-value = 0.0001267
```

```
#SARIMA(1,1,2)x(3,1,0)_12
```

```
model_13_train <- sarima(Avg_ExtentTS_Train, p=1, d=1, q=2, P=3, D=1, Q=0, S=12, details = TRUE)
```


```
## initial value -1.038097  
## iter 2 value -1.174147  
## iter 3 value -1.267295  
## iter 4 value -1.292320  
## iter 5 value -1.298557  
## iter 6 value -1.301947  
## iter 7 value -1.306400  
## iter 8 value -1.321210  
## iter 9 value -1.330951  
## iter 10 value -1.335397  
## iter 11 value -1.337789  
## iter 12 value -1.340572  
## iter 13 value -1.347191  
## iter 14 value -1.348937  
## iter 15 value -1.350213  
## iter 16 value -1.351118  
## iter 17 value -1.351773  
## iter 18 value -1.352098  
## iter 19 value -1.352148
```

```
## iter 20 value -1.352151
## iter 21 value -1.352151
## iter 21 value -1.352151
## final value -1.352151
## converged
## initial value -1.349885
## iter 2 value -1.351590
## iter 3 value -1.352561
## iter 4 value -1.352788
## iter 5 value -1.352968
## iter 6 value -1.352987
## iter 7 value -1.353011
## iter 8 value -1.353019
## iter 9 value -1.353020
## iter 10 value -1.353021
## iter 10 value -1.353021
## iter 10 value -1.353021
## final value -1.353021
## converged
```



```
model_13_train_residuals = resid(model_13_train$fit)
hist(model_13_train_residuals)
```

Histogram of model_13_train_residuals



```
shapiro.test(model_13_train_residuals)
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_13_train_residuals  
## W = 0.98949, p-value = 0.006128
```

```
#SARIMA(2,1,1)x(3,1,0)_12
```


```
model_14_train <- sarima(Avg_ExtentTS_Train, p=2, d=1, q=1, P=3, D=1, Q=0, S=12, details = TRUE)
```

```
## initial value -1.036901  
## iter 2 value -1.166690  
## iter 3 value -1.255867  
## iter 4 value -1.280740  
## iter 5 value -1.285485  
## iter 6 value -1.286803  
## iter 7 value -1.286902  
## iter 8 value -1.286964  
## iter 9 value -1.287912  
## iter 10 value -1.294229  
## iter 11 value -1.298570  
## iter 12 value -1.305005  
## iter 13 value -1.310441  
## iter 14 value -1.314638  
## iter 15 value -1.315596  
## iter 16 value -1.329652  
## iter 17 value -1.330212  
## iter 18 value -1.330363  
## iter 19 value -1.330578
```

```

## iter 20 value -1.330666
## iter 21 value -1.330705
## iter 22 value -1.330707
## iter 23 value -1.330707
## iter 24 value -1.330707
## iter 24 value -1.330707
## iter 24 value -1.330707
## final value -1.330707
## converged
## initial value -1.343351
## iter 2 value -1.348834
## iter 3 value -1.350278
## iter 4 value -1.351392
## iter 5 value -1.351635
## iter 6 value -1.351709
## iter 7 value -1.351762
## iter 8 value -1.351781
## iter 9 value -1.351785
## iter 10 value -1.351785
## iter 10 value -1.351785
## iter 10 value -1.351785
## final value -1.351785
## converged

```




```

model_14_train_residuals = resid(model_14_train$fit)
hist(model_14_train_residuals)

```

Histogram of model_14_train_residuals



```
shapiro.test(model_14_train_residuals)
```

```
##  
##  Shapiro-Wilk normality test  
##  
## data:  model_14_train_residuals  
## W = 0.99044, p-value = 0.01131
```

```
#SARIMA(2,1,2)x(3,1,0)_12
```

```
model_15_train <- sarima(Avg_ExtentTS_Train, p=2, d=1, q=2, P=3, D=1, Q=0, S=12 , details = TRUE)
```

```
## initial  value -1.036901  
## iter    2 value -1.157349  
## iter    3 value -1.257467  
## iter    4 value -1.285105  
## iter    5 value -1.293434  
## iter    6 value -1.295560  
## iter    7 value -1.300881  
## iter    8 value -1.302705  
## iter    9 value -1.311610  
## iter   10 value -1.325367  
## iter   11 value -1.333384  
## iter   12 value -1.335509  
## iter   13 value -1.335519  
## iter   14 value -1.337033  
## iter   15 value -1.338102  
## iter   16 value -1.339632  
## iter   17 value -1.340943  
## iter   18 value -1.341232  
## iter   19 value -1.343535
```


```

## iter 20 value -1.345330
## iter 21 value -1.346789
## iter 22 value -1.346934
## iter 23 value -1.347039
## iter 24 value -1.347063
## iter 25 value -1.347078
## iter 26 value -1.347086
## iter 27 value -1.347112
## iter 28 value -1.347139
## iter 29 value -1.347147
## iter 30 value -1.347149
## iter 31 value -1.347149
## iter 32 value -1.347149
## iter 32 value -1.347149
## iter 32 value -1.347149
## final value -1.347149
## converged
## initial value -1.348512
## iter 2 value -1.349569
## iter 3 value -1.349586
## iter 4 value -1.350848
## iter 5 value -1.350982
## iter 6 value -1.351193
## iter 7 value -1.352121
## iter 8 value -1.352601
## iter 9 value -1.353193
## iter 10 value -1.353272
## iter 11 value -1.353333
## iter 12 value -1.353339
## iter 13 value -1.353341
## iter 14 value -1.353341
## iter 15 value -1.353341
## iter 16 value -1.353342
## iter 16 value -1.353341
## iter 16 value -1.353341
## final value -1.353342
## converged

## Warning in sqrt(diag(fitit$var.coef)): NaNs produced


## Warning in sqrt(diag(fitit$var.coef)): NaNs produced

```

```
model_15_train_residuals = resid(model_15_train$fit)
hist(model_15_train_residuals)
```

Histogram of model_15_train_residuals



```


shapiro.test(model_15_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_15_train_residuals
## W = 0.98876, p-value = 0.003849

#SARIMA(1,1,3)x(3,1,0)_12
model_16_train <- sarima(Avg_ExtentTS_Train, p=1, d=1, q=3, P=3, D=1, Q=0, S=12 , details = TRUE)


## initial  value -1.038097
## iter    2 value -1.179338
## iter    3 value -1.282734
## iter    4 value -1.316080
## iter    5 value -1.323583
## iter    6 value -1.325829
## iter    7 value -1.327176
## iter    8 value -1.330400
## iter    9 value -1.336051
## iter   10 value -1.343009
## iter   11 value -1.345708
## iter   12 value -1.349993
## iter   13 value -1.351352
## iter   14 value -1.351698
## iter   15 value -1.352012
## iter   16 value -1.352157
## iter   17 value -1.352161
## iter   18 value -1.352172
## iter   19 value -1.352179
## iter   20 value -1.352186
## iter   21 value -1.352188
## iter   22 value -1.352189
## iter   23 value -1.352190
## iter   24 value -1.352190
## iter   24 value -1.352190
## iter   24 value -1.352190
## final  value -1.352190
## converged
## initial  value -1.349943
## iter    2 value -1.352612
## iter    3 value -1.352925
## iter    4 value -1.353025
## iter    5 value -1.353075
## iter    6 value -1.353123
## iter    7 value -1.353167
## iter    8 value -1.353192
## iter    9 value -1.353198
## iter   10 value -1.353199
## iter   11 value -1.353199
## iter   11 value -1.353199
## iter   11 value -1.353199
## final  value -1.353199
## converged

```



```
model_16_train_residuals = resid(model_16_train$fit)
hist(model_16_train_residuals)
```

Histogram of model_16_train_residuals



```

shapiro.test(model_16_train_residuals)

##
##  Shapiro-Wilk normality test
##
## data:  model_16_train_residuals
## W = 0.98903, p-value = 0.004551

# This is commented out because it results in an error
#SARIMA(3,1,1)x(3,1,0)_12
#model_17_train <- sarima(Avg_ExtentTS_Train, p=3, d=1, q=1, P=3, D=1, Q=0, S=12, details = TRUE)
#model_17_train_residuals = resid(model_17_train$fit)
#hist(model_17_train_residuals)
#shapiro.test(model_17_train_residuals)
#gives an error when run

#SARIMA(2,1,3)x(3,1,0)_12
model_18_train <- sarima(Avg_ExtentTS_Train, p=2, d=1, q=3, P=3, D=1, Q=0, S=12, details = TRUE)


## initial value -1.036901
## iter 2 value -1.163257
## iter 3 value -1.266869
## iter 4 value -1.296361
## iter 5 value -1.310632
## iter 6 value -1.321697
## iter 7 value -1.325417
## iter 8 value -1.332240
## iter 9 value -1.336089
## iter 10 value -1.340935
## iter 11 value -1.350055
## iter 12 value -1.350337
## iter 13 value -1.350570
## iter 14 value -1.350839
## iter 15 value -1.350916
## iter 16 value -1.351018
## iter 17 value -1.351217
## iter 18 value -1.351605
## iter 19 value -1.351656
## iter 20 value -1.351764
## iter 21 value -1.351810
## iter 22 value -1.351859
## iter 23 value -1.351932
## iter 24 value -1.352073
## iter 25 value -1.352252
## iter 26 value -1.352391
## iter 27 value -1.352416
## iter 28 value -1.352463
## iter 29 value -1.352465
## iter 30 value -1.352465
## iter 31 value -1.352466
## iter 32 value -1.352466
## iter 32 value -1.352466
## iter 33 value -1.352466
## iter 33 value -1.352466
## iter 33 value -1.352466

```

```
## final value -1.352466
## converged
## initial value -1.352034
## iter 2 value -1.352248
## iter 3 value -1.352661
## iter 4 value -1.352988
## iter 5 value -1.353009
## iter 6 value -1.353020
## iter 7 value -1.353021
## iter 8 value -1.353021
## iter 9 value -1.353021
## iter 9 value -1.353021
## iter 9 value -1.353021
## final value -1.353021
## converged


## Warning in sqrt(diag(fitit$var.coef)): NaNs produced

## Warning in sqrt(diag(fitit$var.coef)): NaNs produced
```



```
model_18_train_residuals = resid(model_18_train$fit)
hist(model_18_train_residuals)
```

Histogram of model_18_train_residuals



```
shapiro.test(model_18_train_residuals)
```


```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_18_train_residuals  
## W = 0.9895, p-value = 0.00616
```

```
#SARIMA(3,1,2)x(3,1,0)_12
```

```
model_19_train <- sarima(Avg_ExtentTS_Train, p=3, d=1, q=2, P=3, D=1, Q=0, S=12, details = TRUE)
```


```
## initial value -1.036159  
## iter 2 value -1.160007  
## iter 3 value -1.259848  
## iter 4 value -1.286934  
## iter 5 value -1.296300  
## iter 6 value -1.299039  
## iter 7 value -1.331452  
## iter 8 value -1.341032  
## iter 9 value -1.349800  
## iter 10 value -1.351699  
## iter 11 value -1.352268  
## iter 12 value -1.352842  
## iter 13 value -1.354250  
## iter 14 value -1.357163  
## iter 15 value -1.359003  
## iter 16 value -1.359630  
## iter 17 value -1.360347  
## iter 18 value -1.361172  
## iter 18 value -1.361172
```

```
## iter 18 value -1.361172
## final value -1.361172
## converged
## initial value -1.349476
## iter 2 value -1.350545
## iter 3 value -1.352126
## iter 4 value -1.352892
## iter 5 value -1.353195
## iter 6 value -1.353227
## iter 7 value -1.353242
## iter 8 value -1.353248
## iter 9 value -1.353255
## iter 10 value -1.353260
## iter 11 value -1.353313
## iter 12 value -1.353376
## iter 13 value -1.353392
## iter 14 value -1.353425
## iter 15 value -1.353444
## iter 16 value -1.353444
## iter 17 value -1.353446
## iter 17 value -1.353445
## final value -1.353446
## converged
```



```
model_19_train_residuals = resid(model_19_train$fit)
hist(model_19_train_residuals)
```

Histogram of model_19_train_residuals



```
shapiro.test(model_19_train_residuals)
```

```
##  
## Shapiro-Wilk normality test  
##  
## data:  model_19_train_residuals  
## W = 0.98931, p-value = 0.005438
```

```
#SARIMA(3,1,3)x(3,1,0)_12
```


```
model_20_train <- sarima(Avg_ExtentTS_Train, p=3, d=1, q=3, P=3, D=1, Q=0, S=12, details = TRUE)
```

```
## initial value -1.036159  
## iter 2 value -1.164260  
## iter 3 value -1.261457  
## iter 4 value -1.294471  
## iter 5 value -1.310799  
## iter 6 value -1.331362  
## iter 7 value -1.332748  
## iter 8 value -1.340710  
## iter 9 value -1.342401  
## iter 10 value -1.344768  
## iter 11 value -1.356444  
## iter 12 value -1.358054  
## iter 13 value -1.361687  
## iter 14 value -1.361893  
## iter 15 value -1.361951  
## iter 16 value -1.362041  
## iter 16 value -1.362041  
## iter 17 value -1.362048  
## iter 17 value -1.362048
```




```
## iter 18 value -1.362055
## iter 18 value -1.362055
## iter 19 value -1.362057
## iter 19 value -1.362057
## iter 19 value -1.362057
## final value -1.362057
## converged
## initial value -1.349245
## iter 2 value -1.350423
## iter 3 value -1.351430
## iter 4 value -1.352904
## iter 5 value -1.352954
## iter 6 value -1.353039
## iter 7 value -1.353175
## iter 8 value -1.353367
## iter 9 value -1.353555
## iter 10 value -1.353680
## iter 11 value -1.353702
## iter 12 value -1.353712
## iter 13 value -1.353734
## iter 14 value -1.353767
## iter 15 value -1.353868
## iter 16 value -1.353965
## iter 17 value -1.354067
## iter 18 value -1.354333
## iter 19 value -1.354582
## iter 20 value -1.354746
## iter 21 value -1.354794
## iter 22 value -1.354799
## iter 23 value -1.354846
## iter 24 value -1.354906
## iter 25 value -1.355108
## iter 26 value -1.356058
## iter 27 value -1.356158
## iter 28 value -1.356195
## iter 29 value -1.357452
## iter 30 value -1.358237
## iter 31 value -1.358849
## iter 32 value -1.359578
## iter 33 value -1.360331
## iter 34 value -1.361291
## iter 35 value -1.361471
## iter 36 value -1.362078
## iter 37 value -1.362381
## iter 38 value -1.362544
## iter 39 value -1.362731
## iter 40 value -1.362773
## iter 41 value -1.362816
## iter 42 value -1.362843
## iter 43 value -1.362863
## iter 44 value -1.362906
## iter 45 value -1.362946
## iter 46 value -1.363173
## iter 47 value -1.363342
```

```
## iter 48 value -1.363621
## iter 49 value -1.363728
## iter 50 value -1.363844
## iter 51 value -1.363959
## iter 52 value -1.364065
## iter 53 value -1.364278
## iter 54 value -1.364305
## iter 55 value -1.364463
## iter 56 value -1.364505
## iter 57 value -1.364535
## iter 58 value -1.364554
## iter 59 value -1.364565
## iter 60 value -1.364576
## iter 61 value -1.364584
## iter 62 value -1.364587
## iter 63 value -1.364587
## iter 64 value -1.364587
## iter 64 value -1.364587
## final value -1.364587
## converged
```



```
model_20_train_residuals = resid(model_20_train$fit)
hist(model_20_train_residuals)
```

Histogram of model_20_train_residuals



```
shapiro.test(model_20_train_residuals)
```

```
##
## Shapiro-Wilk normality test
##
## data:  model_20_train_residuals
## W = 0.99038, p-value = 0.01088
```

Summarize the fits of these models in a table.

```
library(huxtable)
goodness_of_fit <- hux(
  Model = c('SARIMA(0,1,2)x(0,1,1)_12', 'SARIMA(0,1,2)x(3,1,0)_12', 'SARIMA(0,1,2)x(1,1,1)_12',
    'SARIMA(4,1,0)x(0,1,1)_12', 'SARIMA(4,1,0)x(3,1,0)_12', 'SARIMA(4,1,0)x(1,1,1)_12',
    'SARIMA(5,1,0)x(0,1,1)_12', 'SARIMA(5,1,0)x(3,1,0)_12', 'SARIMA(5,1,0)x(1,1,1)_12',
    'SARIMA(1,1,1)x(0,1,1)_12', 'SARIMA(1,1,1)x(3,1,0)_12', 'SARIMA(1,1,1)x(1,1,1)_12',
    'SARIMA(1,1,2)x(3,1,0)_12', 'SARIMA(2,1,1)x(3,1,0)_12', 'SARIMA(2,1,2)x(3,1,0)_12',
    'SARIMA(1,1,3)x(3,1,0)_12', 'SARIMA(2,1,3)x(3,1,0)_12', 'SARIMA(3,1,2)x(3,1,0)_12',
    'SARIMA(3,1,3)x(3,1,0)_12'),
  AIC = c(model_1_train$AIC, model_2_train$AIC, model_3_train$AIC,
    model_4_train$AIC, model_5_train$AIC, model_6_train$AIC,
    model_7_train$AIC, model_8_train$AIC, model_9_train$AIC,
    model_10_train$AIC, model_11_train$AIC, model_12_train$AIC,
    model_13_train$AIC, model_14_train$AIC, model_15_train$AIC,
    model_16_train$AIC, model_18_train$AIC, model_19_train$AIC,
    model_20_train$AIC),
  AICc = c(model_1_train$AICc, model_2_train$AICc, model_3_train$AICc,
    model_4_train$AICc, model_5_train$AICc, model_6_train$AICc,
    model_7_train$AICc, model_8_train$AICc, model_9_train$AICc,
    model_10_train$AICc, model_11_train$AICc, model_12_train$AICc,
    model_13_train$AICc, model_14_train$AICc, model_15_train$AICc,
```

```

        model_16_train$AICc, model_18_train$AICc, model_19_train$AICc,
        model_20_train$AICc),
  BIC = c(model_1_train$BIC, model_2_train$BIC, model_3_train$BIC,
    model_4_train$BIC, model_5_train$BIC, model_6_train$BIC,
    model_7_train$BIC, model_8_train$BIC, model_9_train$BIC,
    model_10_train$BIC, model_11_train$BIC, model_12_train$BIC,
    model_13_train$BIC, model_14_train$BIC, model_15_train$BIC,
    model_16_train$BIC, model_18_train$BIC, model_19_train$BIC,
    model_20_train$BIC),
  MSE = c(mean(model_1_train_residuals^2), mean(model_2_train_residuals^2), mean(model_3_train_re
    mean(model_4_train_residuals^2), mean(model_5_train_residuals^2), mean(model_6_train_re
    mean(model_7_train_residuals^2), mean(model_8_train_residuals^2), mean(model_9_train_re
    mean(model_10_train_residuals^2), mean(model_11_train_residuals^2), mean(model_12_train
    mean(model_13_train_residuals^2), mean(model_14_train_residuals^2), mean(model_15_train
    mean(model_16_train_residuals^2), mean(model_18_train_residuals^2), mean(model_19_train
    mean(model_20_train_residuals^2))
)

goodness_of_fit %>%
  set_number_format(col=c(2,3,4,5), value=3) %>%
  set_bottom_border(c(1,13,16), everywhere) %>%
  set_bold(c(12,14,15,16), everywhere) %>%
  set_background_color(evens, everywhere, "grey95")


```

Model	AIC	AICc	BIC	MSE
SARIMA(0,1,2)x(0,1,1)_12	0.260	0.260	0.302	0.070
SARIMA(0,1,2)x(3,1,0)_12	0.274	0.275	0.336	0.071
SARIMA(0,1,2)x(1,1,1)_12	0.264	0.264	0.315	0.070
SARIMA(4,1,0)x(0,1,1)_12	0.258	0.258	0.320	0.069
SARIMA(4,1,0)x(3,1,0)_12	0.275	0.276	0.358	0.070
SARIMA(4,1,0)x(1,1,1)_12	0.261	0.262	0.333	0.069
SARIMA(5,1,0)x(0,1,1)_12	0.247	0.248	0.320	0.068
SARIMA(5,1,0)x(3,1,0)_12	0.264	0.265	0.357	0.069
SARIMA(5,1,0)x(1,1,1)_12	0.251	0.252	0.333	0.068
SARIMA(1,1,1)x(0,1,1)_12	0.305	0.305	0.346	0.073
SARIMA(1,1,1)x(3,1,0)_12	0.195	0.195	0.257	0.064
SARIMA(1,1,1)x(1,1,1)_12	0.308	0.309	0.360	0.069
SARIMA(1,1,2)x(3,1,0)_12	0.168	0.169	0.241	0.062
SARIMA(2,1,1)x(3,1,0)_12	0.171	0.171	0.243	0.062
SARIMA(2,1,2)x(3,1,0)_12	0.173	0.174	0.255	0.062
SARIMA(1,1,3)x(3,1,0)_12	0.173	0.174	0.256	0.062
SARIMA(2,1,3)x(3,1,0)_12	0.179	0.180	0.272	0.062
SARIMA(3,1,2)x(3,1,0)_12	0.178	0.179	0.271	0.062
SARIMA(3,1,3)x(3,1,0)_12	0.161	0.162	0.264	0.060


Model Selection

We evaluate performance on the test set of a few of the models which gave the best fit.


```
model_11_train_forecast<- sarima.for(Avg_ExtentTS_Train, n.ahead=22, p=1,d=1,q=1,P=3,D=1,Q=0,S=12)
lines(Avg_ExtentTS_Test,col='blue',type='b',pch=16)
```




```
model_13_train_forecast<- sarima.for(Avg_ExtentTS_Train, n.ahead=22, p=1,d=1,q=2,P=3,D=1,Q=0,S=12)
lines(Avg_ExtentTS_Test,col='blue',type='b',pch=16)
```



```
model_14_train_forecast<- sarima.for(Avg_ExtentTS_Train, n.ahead=22, p=2,d=1,q=1,P=3,D=1,Q=0,S=12)
lines(Avg_ExtentTS_Test,col='blue',type='b',pch=16)
```



```
model_15_train_forecast<- sarima.for(Avg_ExtentTS_Train, n.ahead=22, p=2,d=1,q=2,P=3,D=1,Q=0,S=12)
lines(Avg_ExtentTS_Test,col='blue',type='b',pch=16)
```



```
mean((model_11_train_forecast$pred-Avg_ExtentTS_Test)^2)
```

```
## [1] 0.3321616
```

```
mean((model_13_train_forecast$pred-Avg_ExtentTS_Test)^2)
```

```
## [1] 0.3385571
```

```
mean((model_14_train_forecast$pred-Avg_ExtentTS_Test)^2)
```

```
## [1] 0.3363344
```

```
mean((model_15_train_forecast$pred-Avg_ExtentTS_Test)^2)
```

```
## [1] 0.3388543
```

Summarize these results in a table.

```
sarima_prediction <- hux(
  Model = c('SARIMA(1,1,1)x(3,1,0)_12', 'SARIMA(1,1,2)x(3,1,0)_12', 'SARIMA(2,1,1)x(3,1,0)_12', 'SARIMA(2,1,2)x(3,1,0)_12'),
  PMSE = c(mean((model_11_train_forecast$pred-Avg_ExtentTS_Test)^2), mean((model_13_train_forecast$pred-Avg_ExtentTS_Test)^2),
            mean((model_14_train_forecast$pred-Avg_ExtentTS_Test)^2), mean((model_15_train_forecast$pred-Avg_ExtentTS_Test)^2))

sarima_prediction %>%
  set_number_format(col=2, value=3) %>%
  set_bottom_border(1, everywhere) %>%
  set_background_color(evens, everywhere, "grey95")
```


Model	PMSE
SARIMA(1,1,1)x(3,1,0)_12	0.332
SARIMA(1,1,2)x(3,1,0)_12	0.339
SARIMA(2,1,1)x(3,1,0)_12	0.336
SARIMA(2,1,2)x(3,1,0)_12	0.339